From Gates to a RISC-V CPU

Hands-on digital design on an FPGA,
with an open-source toolchain

Lattice ECP5 `LFE5U-85F` · Radiona ULX3S
Yosys · nextpnr · Project Trellis · Icarus Verilog · Verilator

Version 1.0 · 11 November 2026

This book is licensed under Creative Commons Attribution 4.0 International.
The accompanying code is licensed under the MIT License.
Copyright © 2026 The Open FPGA Course.

Contents

About this course 1
- What you will build 1
- The method 2
- Two ways to do it 2
- What this course is not 2

The ULX3S, and what is inside it 3
- The board 3
- The chip 3
- What is actually inside 3
- Why an open toolchain matters here 4

The toolchain, and the people who built it 5
- The chain from Verilog to bitstream 5
- The board 5
- The instruction set 5
- The visual path 6
- On standing on shoulders 6
- This course's own licence 6

Installation 7
- The short version 7
- What the installer does 7
- Verifying it 7
- The board 8
- Editor setup 8
- Icestudio, for the visual path 8

How to run the exercises 9
- The shape of a lesson 9
- The commands 9
- The order to do things in 9
- On the exercises 10
- When something goes wrong 10
- Keeping the whole course green 10

Lesson 00 — Blinky 11
- The idea 11
- Theory 11
- What it cost 12
- Learning outcomes 12

Lesson 01 — Logic gates 13
- The idea 13
- Theory 13
- What it cost, and why it is the point 13
- Learning outcomes 14

Lesson 02 — Combinational logic 15
- The idea 15
- Theory 15
- What it cost 16
- Learning outcomes 16

Lesson 03 — Sequential logic 17
- The idea 17
- Theory 17
- What it cost 18
- Learning outcomes 19

Lesson 04 — Clocks and timing 20
The idea 20
The PLL is a hard block 20
Crossing between domains 21
What was measured 21
The honest part about forcing a failure 21
Lesson 05 — Finite state machines 23
The idea 23
Moore and Mealy 23
Two textbook rules, measured 23
The mistake everyone makes 24
Testing changes shape 24
What it cost 24
Lesson 06 — SPI and the OLED 25
The idea 25
Theory 25
What it cost 26
Testing an interface 26
Lesson 07 — Font ROM and coordinate maths 27
The idea 27
Theory: make the arithmetic disappear 27
What it cost, and the question it raises 28
On what the test does not check 28
Lesson 08 — Memories 29
The idea 29
The measurement 29
Lesson 07's puzzle, answered 30
Two traps found while building it 30
Why a framebuffer wants two ports 30
Lesson 09 — UART 31
The idea 31
Where the textbook had to be overruled 31
The honest limit 32
What a loopback test does not prove 32
What it cost 32
Lesson 10 — The ALU 34
The idea 34
Theory 34
Exhaustive testing dies here 35
The surprise in the resource report 35
Timing: the ceiling lesson 13 will hit 35
Lesson 11 — The register file 36
The shape is dictated by the instruction set 36
`x0` 36
Combinational reads, and why lesson 08 agrees 36
Read-during-write returns the old value 37
A mutation that should survive 37
Exhaustive testing is briefly back 37
Lesson 12 — Decode and control 38
Fixed fields, decoded in parallel 38
The immediates look like vandalism 38
The trap 39
Testing a decoder 39

What it cost ... 40
Lesson 13 — A single-cycle RV32I processor ... 41
Nothing here is new ... 41
Why single-cycle first ... 41
The new parts ... 42
The bug only running software could find ... 42
Testing through the architectural interface ... 42
What it cost, and what it cost in time ... 42
On the board ... 43
Lesson 14 — Formal verification ... 44
Why this lesson exists here ... 44
The idea ... 44
`riscv-formal` ... 44
What it actually did ... 44
Catching a real bug ... 45
The failure mode that makes formal dangerous ... 45
What this lesson does not do ... 46
Lesson 15 — A system on chip ... 47
Where a CPU becomes a computer ... 47
The toolchain path ... 47
The payoff ... 47
Lesson 16 — Pipelining ... 48
The argument, already made ... 48
The cost: hazards ... 48
What to measure ... 48
Lesson 17 — Extensions ... 50
Optional, and genuinely so ... 50
Traps ... 50
The machine timer ... 50
The M extension ... 50
Appendix A — Board pinout ... 51
Clock ... 51
LEDs ... 51
Buttons ... 51
Slide switches ... 51
OLED — SSD1331, 96 × 64, 16-bit colour ... 51
USB serial — FTDI FT231X ... 51
Everything else ... 52
Part facts worth remembering ... 52
Appendix B — Command reference ... 53
Per-lesson verbs ... 53
From the project root ... 53
Driving the tools directly ... 53
Reading a resource report ... 54
Which simulator when ... 54
Tools written for this course ... 54
Appendix C — Verilog style rules ... 55
Synchronous design, one clock ... 55
Blocking versus non-blocking ... 55
Do not infer latches ... 55
Widths ... 56
Reset and initialisation ... 56
One driver per signal ... 56

Naming 56
Testbenches 56
Mutation testing 57
Appendix D — Where to go next 58
Specifications worth having open 58
Books 58
Projects to read 58
Things this course deliberately left out 58
A closing thought 59

About this course

This is a hands-on refresher in digital design, worked from a single logic gate to a 32-bit RISC-V processor running compiled C, with a toolchain that is open source from the Verilog parser to the bitstream.

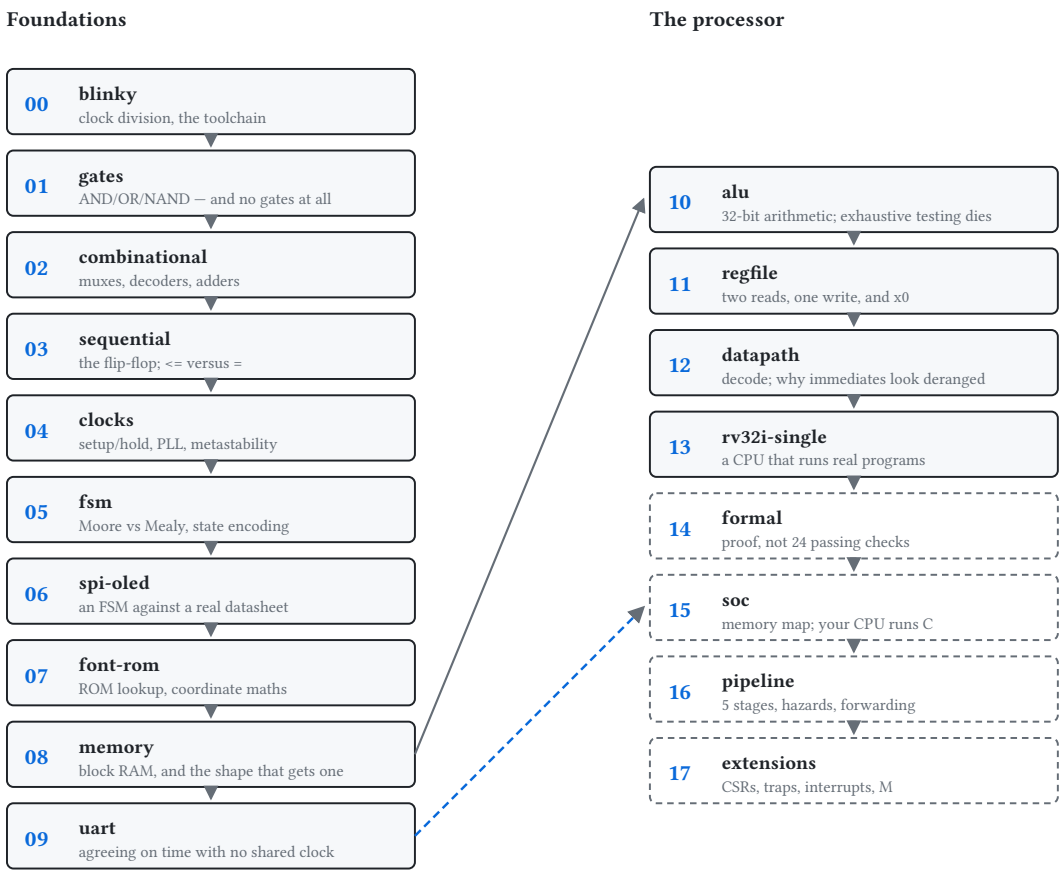
The lessons are built and measured on a **Radiona ULX3S**, carrying a Lattice ECP5 `LFE5U-85F`, and every number in this book was measured on that part. But roughly half the course is board-independent: gates, multiplexers, state machines, the ALU, the register file and instruction decode would be identical on any FPGA. Where a result depends on this particular fabric — and several of the most interesting ones do — the book says so.

It is written for someone who studied this once — perhaps a long time ago, as an electrical engineer — and wants the ideas back **in their hands** rather than back in their head. So concepts are stated briefly and then built: simulated, inspected in a waveform viewer, synthesised, measured, and run on real silicon.

It assumes you know what a gate, a register and a clock are. It does not assume you remember what a blocking assignment does to a shift register, or why your counter synthesised into a latch.

What you will build

Eighteen lessons, each self-contained, each ending in something you can see working. The sequence is not arbitrary — every lesson exists because a later one needs it.



Solid boxes are written and verified. Dashed boxes have settled scope and are being built in order.
The blue arrow is lesson 09's UART becoming the processor's console in lesson 15.

The course. Solid arrows are prerequisites; the right-hand column is the processor, assembled from parts built in the left-hand ones.

By lesson 13 you have a RISC-V CPU running programs. By lesson 15 it runs C compiled by a standard GNU toolchain, printing to your terminal. Lesson 16 pipelines it.

The method

Three habits run through every lesson, and they matter more than any particular circuit.

Simulate first, and make the testbench say PASS or FAIL

Nothing goes near synthesis until a self-checking testbench passes. Every testbench in this course prints an explicit verdict — never a waveform to squint at — because a test you have to interpret is a test you will misinterpret at one in the morning.

Measure; do not inherit opinions

This is the one that makes the course worth doing rather than reading. The synthesiser's output is two seconds away, and it repeatedly contradicts what everyone **knows**:

> **MEASURED four received rules that did not survive contact with the fabric**
>
> - "One-hot encoding suits FPGAs, since flip-flops are free." On this four-state machine it cost **41 LUT4 against 12** — nearly four times worse (lesson 05).
> - "Mealy machines need fewer states, so they are cheaper." Fewer states, yes — 3 against 4 — and **more logic**: 5 LUT4 against 2 (lesson 05).
> - "The synthesiser will reduce a constant multiply to shifts." It spent **two hardware multiplier blocks** instead (lesson 08).
> - "Build it from gates to understand it, then it is the same thing." The hand-built 32-bit adder cost **212 cells to +'s 17** (lesson 02).

None of those are in this book because they were read somewhere. They are here because `make synth` said so, and the numbers are reproducible on your machine in the time it takes to read this paragraph.

Account for every flip-flop

When `make synth` reports more state than you can explain, you have inferred something you did not intend. Lesson 05's design reports 113 flip-flops, and they decompose exactly: 25 for the tick divider, 5 for the traffic light, 76 for four debouncers, 4 for edge detectors, 3 for the lock's state. That habit finds bugs no test was written for.

Two ways to do it

Every lesson up to 09 can be done by typing Verilog in an editor, or by wiring blocks together visually in Icestudio. Both are first class, and you can switch at any lesson.

The two cannot drift apart, because the Icestudio project is *generated* from the lesson's real Verilog: a short specification names the source file and binds each port to a board pin, and a tool lifts the module body into a visual code block. The diagram you wire up and the file you would type are the same design, laid out differently.

Going the other way, `make graph` runs your Verilog through the synthesiser and hands you back a schematic. Doing that once per lesson is the fastest way to find the gap between what you meant and what you said.

> **WHY THIS MATTERS**
>
> Beyond lesson 09 the lessons are editor-only, and that is itself a lesson. Schematic entry stops scaling at roughly the complexity of a 32-bit datapath: a 32-bit ALU drawn as blocks is a wall of identical slices nobody can read, while as text it is a dozen lines with `[31:0]` on them. The industry made the same transition in the late 1980s for the same reason.

What this course is not

It is not a Verilog language reference, and it is not a complete treatment of computer architecture. Both exist and are better at being themselves; several are listed at the end of this book.

It is also not a tour of IP cores. Nothing here is instantiated from a vendor library. Every component — the adder, the register file, the UART, the CPU — is built and tested in a lesson, because the point is to have built it.

The ULX3S, and what is inside it

The board

The ULX3S is an open-hardware FPGA development board from Radiona, the makerspace in Zagreb. Schematics, PCB layout and a large body of example gateware are published openly at github.com/emard/ulx3s.

It was designed for teaching, and it shows: the peripherals are the ones you actually want when learning, placed where you can reach them, with no mezzanine cards to buy.

Clock	25 MHz, free-running, straight from a crystal
LEDs	8, active high
Buttons	7, undebounced — which lesson 03 turns into a virtue
Switches	4 slide switches, static
Display	SSD1331 colour OLED, 96×64, 16-bit colour, 4-wire SPI
Serial	FTDI FT231X — also the JTAG programmer, over one USB cable
Memory	32 MB SDRAM, 16 bits wide; SPI flash for the bitstream
Storage	micro-SD slot
Video	digital video output
Expansion	28 differential GPIO pairs on two headers
Wireless	an ESP32 module

Four of those matter for this course: the clock, the LEDs, the buttons and the OLED. Lessons 00–13 need no more. The SDRAM, the SD card and video output are where you go after the course ends.

The chip

The board in this course carries a Lattice `LFE5U-85F-6BG381C`. Unpacking that part number is worth a paragraph, because every field of it turns up later:

`LFE5U`	Lattice ECP5, the "U" being the version without SERDES
`85`	the size grade — about 84,000 logic cells
`F`	the package family
`6`	speed grade 6, the middle of three. A faster part would close timing at a higher clock with identical Verilog
`BG381`	a 381-ball caBGA
`C`	commercial temperature range

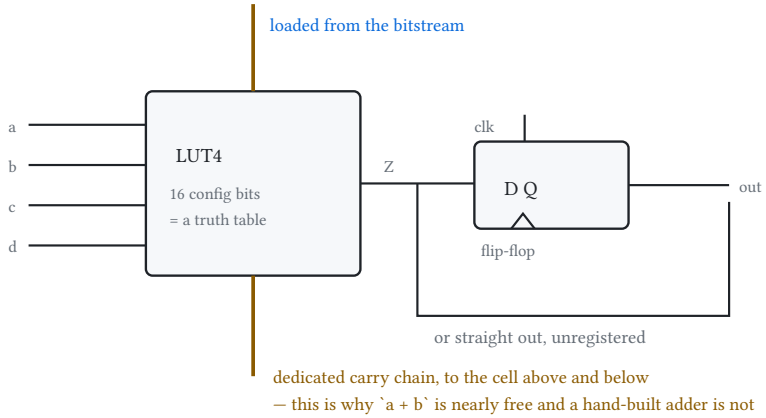
What you get:

resource	count	what it is
LUT4	83,640	4-input lookup tables — the general-purpose logic
Flip-flops	83,640	one per logic cell, whether you use it or not
Block RAM	208 × 18 Kbit	about 3.7 Mbit of dedicated memory
Multipliers	156	18×18 hardware multipliers
PLLs	4	frequency synthesis, lesson 04

For calibration: lesson 13's whole RISC-V CPU uses about 1,765 LUT4 — a little over **2%** of the part. Logic capacity will not be your constraint in this course, or in most projects.

What is actually inside

This is the single most useful piece of theory in the book, because almost every cost surprise in the course follows from it.



Every logic cell has a flip-flop whether the design uses it or not, which is why flip-flops are not the scarce resource.
Any function of up to four inputs costs exactly one LUT4, however complicated: a 2-input AND and a 4-input XOR are the same price.

An ECP5 logic cell. The LUT's sixteen configuration bits come from the bitstream, which is what "programming" an FPGA means.

An FPGA contains no gates. It contains **lookup tables**: sixteen configuration bits holding the truth table of whatever four-input function you asked for, selected by the four inputs. Your `a & b` does not become an AND gate; it becomes a LUT4 loaded with the AND truth table and two inputs unused.

Three consequences run through the whole course:

1. **Any function of up to four inputs costs exactly one LUT4**, however complicated. A 2-input AND and a 4-input XOR are the same price (lesson 01). Minimising gate count — half of a classic logic design course — mostly does not apply here.
2. **Flip-flops are not scarce.** Every logic cell has one whether the design uses it or not. This is why one-hot state encoding is **supposed** to be free on an FPGA — and lesson 05 measures a case where it is four times worse anyway.
3. **Some things have dedicated hardware and some do not.** There is a carry chain running up each column of cells, so addition is nearly free: a 32-bit adder costs 17 carry cells and no LUTs. There is no dedicated shift hardware, so a 32-bit barrel shifter costs about 667 LUT4 — which is why lesson 10 finds that the shifter, not the adder, is 80% of the ALU.

> **WHY THIS MATTERS**
> "Which operations are expensive" is not intuition. It is a property of the fabric in front of you, and on a chip without a carry chain the answer would invert. This is why the course measures everything rather than citing rules.

Why an open toolchain matters here

The ECP5 can be programmed by entirely open-source software, from Verilog source to a bitstream the chip accepts. That is unusual, and it is not an accident of licensing — it is the result of years of work reverse-engineering and documenting the bitstream format, described in the next chapter.

It matters for a course for three reasons:

- **Nothing expires and nothing phones home.** No licence server, no registration, no node-locked key that stops working when you change laptop.
- **You can read the tools.** When yosys does something surprising — and this course documents several cases — the source is there. The schematic `make graph` produces is the tool showing you its own intermediate representation.
- **It runs anywhere.** The entire toolchain for this course is one download of about a gigabyte, needs no administrator rights, and is removed by deleting a directory.

The toolchain, and the people who built it

None of this course would exist without a long chain of mostly volunteer work, much of it given away. The tools are named throughout the lessons; this chapter is where the people are.

The chain from Verilog to bitstream

Icarus Verilog	Simulation. The `make sim` in every lesson. Written and maintained by **Stephen Williams** since 1998 — one of the oldest free Verilog implementations, and still the most forgiving one to learn in.
Verilator	The strict linter, and a very fast simulator. Created by **Wilson Snyder**. Every `make lint` in this course is Verilator, and it caught several real bugs during the writing — including a register driven from two `always` blocks and a width truncation that a passing test had hidden.
Yosys	Synthesis: Verilog in, a netlist of ECP5 primitives out. Created by **Claire Wolf**, now developed at YosysHQ. Every resource measurement in this book is `yosys … stat`, and every schematic is `yosys show`.
nextpnr	Place and route, and timing analysis. Created by **gatecat** (formerly David Shah) with contributions from many others. Every "maximum frequency" figure in this book comes from nextpnr.
Project Trellis	The ECP5 bitstream documentation — the work that makes all of the above possible on this chip. Also **gatecat**'s, with contributors. `ecppack` and the PLL calculator `ecppll` are from this project.
GTKWave	Waveform viewing, by **Tony Bybell**. The `make wave` of every lesson.
openFPGALoader	Getting the bitstream onto the board, across a remarkable range of cables and boards. By **Gwenhael Goavec-Merou**.
SymbiYosys	Formal verification, used in lesson 14. From YosysHQ.
OSS CAD Suite	All of the above, built nightly for every platform as one archive, by YosysHQ. `tools/install.sh` downloads exactly this, which is why setting up this course takes one command instead of an afternoon.

> **WHY THIS MATTERS**
> Project IceStorm — Claire Wolf's documentation of the Lattice iCE40 bitstream, published in 2015 — is the root of this whole tree. It demonstrated that a modern FPGA's bitstream could be understood well enough to target with free tools, and Project Trellis then did the same for the much larger ECP5. Everything in this course descends from that demonstration.

The board

ULX3S	Open hardware from **Radiona.org**, the Zagreb makerspace, developed with **Goran Mahovlić** and **Davor Jadrijević** ("emard"), who also maintains an extensive collection of example designs — SDRAM controllers, video output, emulators of older machines — that is worth reading once this course is finished.
Pin constraints	`common/ulx3s_v20.lpf` in this repository comes from the ULX3S project itself. It is the one file here that is not original work, and it is the file that makes everything else possible: it is the map from names like `led[0]` to balls like `B2`.

The instruction set

RISC-V	Designed at the University of California, Berkeley, by **Krste Asanović**, **Andrew Waterman**, **Yunsup Lee** and **David Patterson**, and released as an open standard. The specification is free to read and free to implement, which is why a course can end with "and now build the processor" rather than "and now read about one".
riscv-formal	A formal specification of RV32I that an implementation can be checked against, by **Claire Wolf**. Lesson 14 uses it to prove the lesson 13 core correct — not merely tested.

The quality of the RISC-V specification is itself a teaching aid. Lesson 12 spends several pages on **why** the immediate fields are scattered the way they are, and the answer — that the layout minimises multiplexer inputs in the decoder, with 24 of 32 immediate bits needing no multiplexer at all — is a designed property you can measure, not a historical accident.

The visual path

Icestudio	The block-diagram editor used as this course's alternative path, by **Jesús Arroyo Torrens** and the **FPGAwars** community.
apio	The build system underneath Icestudio, also from FPGAwars, which wraps the same yosys and nextpnr that the command line uses. The two paths in this course are two front ends onto one toolchain, not two different worlds.

On standing on shoulders

It is worth being explicit about the scale of what is being given away here. A commercial FPGA toolchain is a multi-hundred-megabyte, licence-managed install from a vendor with a commercial interest in the chips it targets. The equivalent used in this course was largely written by individuals, is readable end to end, and asks nothing.

If this course is useful to you, the projects above are where the gratitude should go — and several of them take contributions, bug reports and documentation from people whose first FPGA design was a blinking LED not long before.

This course's own licence

In the same spirit, and chosen so that reuse is easy rather than merely permitted:

Code	MIT. The Verilog, the tools, the Makefiles and the assembly programs. Use them in anything, commercial or not; keep the copyright notice.
Prose	Creative Commons Attribution 4.0. The lessons, this book and its diagrams. Share them, adapt them, translate them, teach from them — commercially if you like — giving credit and indicating changes.

One file is not this project's to license: `common/ulx3s_v20.lpf`, the board's pin constraints, which come from the ULX3S project and are included because the course cannot work without them. Everything else — every diagram, every line of Verilog, and the whole of this book — is original to the course.

Installation

Everything needed for lessons 00 through 14 is one command and about a gigabyte. Nothing is installed system-wide, nothing needs administrator rights, and removing the `toolchain/` directory undoes all of it.

The short version

```
git clone <this repository> && cd ULX3S
tools/install.sh
make check
```

`make check` verifies that every tool actually runs and then simulates lesson 00 as a smoke test. It should end with `Ready.`

What the installer does

`tools/install.sh` detects your platform, asks GitHub for the newest OSS CAD Suite release rather than pinning a date that rots, downloads it into `toolchain/`, and — on macOS — clears the quarantine flags that would otherwise make every tool prompt for confirmation.

flag	adds	needed by
none	yosys, nextpnr-ecp5, Project Trellis, Icarus Verilog, Verilator, GTKWave, openFPGALoader, SymbiYosys	lessons 00–14
`--with-riscv`	a RISC-V bare-metal C compiler	lesson 15 onward
`--with-docs`	mdBook, for the HTML version of the lessons	optional
`--all`	all of the above	

Supported platforms are macOS (Intel and Apple Silicon) and Linux (x86-64 and ARM64). On an unsupported platform the script says so and tells you where to put a toolchain you build yourself.

Verifying it

```
make check
```

✓	`iverilog`	simulation
✓	`yosys`	synthesis
✓	`nextpnr-ecp5`	place and route
✓	`ecppack`	bitstream packing
✓	`verilator`	linting
✓	`gtkwave`	waveforms
✓	`dot`	schematics
✓	`make sim`	`PASS: 8 toggles, period 256 clks`

A missing **optional** tool is reported with a dash rather than a cross, and the check still passes — `openFPGALoader` is only needed with a board attached, `sby` only in lesson 14, a RISC-V compiler only from lesson 15.

The board

The ULX3S appears as an FTDI device, vendor `0403`, product `6015`. One USB cable carries both the programmer and the serial port.

```
toolchain/oss-cad-suite/bin/openFPGALoader --detect -b ulx3s
```

A working board answers with its JTAG identity:

```
idcode 0x41113043
manufacturer lattice
family ECP5
model  LFE5U-85
```

On macOS this works with no driver installation. On Linux you need udev rules so that programming does not require root; the rules file shipped with openFPGALoader is the one to use, followed by `sudo udevadm control --reload`.

> **WATCH OUT**
> The ULX3S has two USB connectors. Only one of them is wired to the FPGA's programmer. If `--detect` finds nothing, try the other one before suspecting anything else — and check the cable is a data cable rather than a charge-only one, which is the other common cause of a board that appears dead.

Editor setup

Any editor works, because the Makefiles are the interface. For VS Code the repository ships a complete configuration:

```
tools/code.sh
```

Launch it that way rather than from the Dock. Editor extensions inherit the environment of whatever started the editor, and from a desktop launcher that environment has a bare `PATH` — so the Verilog extension cannot find `verilator` and linting silently does nothing. `tools/code.sh` puts the toolchain on `PATH` first.

The workspace then offers one task per lesson, so running a lesson is **Terminal → Run Task → ▶ `deploy 00 blinky`** — two clicks, and it does not matter which file is open.

> **WATCH OUT**
> VS Code tasks can only read the **focused editor tab**; there is no way for them to see what is highlighted in the file tree. So clicking a folder in the sidebar genuinely cannot select a lesson, no matter how the tasks are configured. That is why the repository enumerates the lessons as tasks instead.

Icestudio, for the visual path

Icestudio is a separate download from icestudio.io. Set the board to **ULX3S-85F** once, in the setup wizard, and the lessons' `.ice` projects open and build with no further configuration.

Icestudio brings its own copy of the toolchain, so the two paths do not share `toolchain/` — which is harmless, if a little redundant on disk.

How to run the exercises

The shape of a lesson

Every lesson is a directory holding its own source, testbench, Makefile and written lesson:

```
13-rv32i-single/
  README.md        the lesson
  cpu.v            synthesisable source
  tb_cpu.v         testbench — always prints PASS or FAIL
  test.s           an assembly program, where the lesson needs one
  Makefile         four lines; the rules live in common/lesson.mk
  diagrams/        timing diagram sources
  icestudio/       the visual path, where the lesson has one
```

Each one follows the same arc: the idea, the design, simulate it, look at the waveform, see what it cost, run it on the board, exercises.

The commands

Every lesson understands the same verbs. Run them from inside a lesson folder:

`make sim`	Icarus Verilog — run the testbench. **Must print PASS.**
`make wave`	GTKWave — look at the waveform
`make lint`	Verilator `-Wall` — latches, widths, multiple drivers
`make graph`	yosys + graphviz — render the design as a schematic
`make synth`	the resource report: LUT4s, flip-flops, block RAM, carry cells
`make bit`	yosys → nextpnr → ecppack — a bitstream, with timing closure
`make prog`	load into FPGA SRAM — volatile, for iterating
`make flash`	write to SPI flash — survives power-off
`make clean`	delete build products

Or from the project root, without changing directory — `L` is the lesson number:

```
make list              # which lesson is which number
make sim    L=5        # simulate lesson 05
make deploy L=0        # build lesson 00 and load it onto the board
make wave   L=3        # open lesson 03's waveform
```

And across the whole course:

`make check`	is the toolchain healthy?
`make sim-all`	every lesson's testbench — the regression suite
`make lint-all`	Verilator across every lesson
`make ice`	regenerate the Icestudio projects from their specs
`make docs`	build the HTML version of the lessons
`make pdf`	build this book

The order to do things in

1. **Read the lesson's "The idea" section** before looking at any code. Most lessons are about one idea and the code is its consequence.
2. **`make sim`.** It must print `PASS`. If it does not, the lesson is not finished and nothing else is worth doing yet.
3. **`make wave`**, and look at the three signals the lesson tells you to look at. A waveform with forty signals in it is not useful until someone tells you which three matter, which is what the lessons do.

4. **`make synth`**, and read the number. Then try to account for it. This is the habit that separates "it works" from "I know what I built".
5. **`make graph`**, at least once per lesson. Seeing what the synthesiser inferred from what you wrote is the fastest way to find the gap between the two.
6. **`make prog`**, where the lesson has a board demo.
7. **The exercises.** Most of them break something deliberately. Predict what will happen before you run it — a prediction you can check is worth ten instructions you can follow.

On the exercises

They are not homework. Each is chosen because it teaches something by failing in an instructive way, and each states what you should observe when it works.

Three kinds recur:

- **Break it and watch.** "Change `⩽` to `=` and run `make sim`." The failure is the lesson.
- **Predict then measure.** "How many flip-flops will this use? Now run `make synth`."
- **Extend it.** "Add a parity bit." These are the ones that take an evening, and the ones worth doing.

Worked solutions are shipped only where an answer is genuinely hard to self-check — lesson 01 has a set; most lessons do not need one, because the exercise states what you should see.

When something goes wrong

`command not found`	the toolchain is missing — run `tools/install.sh`
macOS: "cannot be opened"	quarantine flags; `xattr -dr com.apple.quarantine toolchain/oss-cad-suite`
linting does nothing	the editor was launched without the toolchain on `PATH` — use `tools/code.sh`
`No rule to make target 'bit'`	`make` ran in the project root with no lesson named. Add `L=n`, or open a file inside a lesson first
`Nothing to be done for 'bit'`	the bitstream is newer than the source. `make clean` to force a rebuild
`openFPGALoader` finds nothing	wrong USB socket, or a charge-only cable
"no extension found for verilog"	you pressed the editor's Run/Debug button. Verilog is not **executed** — it is simulated or synthesised. Use `make sim`
`nextpnr` fails timing	the design is too slow for the requested clock. Read the critical path it reports; lesson 04 is about exactly this

Keeping the whole course green

```
make sim-all
```

A later lesson must never break an earlier one, so the whole course is a regression suite. When a module is promoted into `common/rtl/` — as the adder and the debouncer are, once a second lesson needs them — this is what catches a change that breaks a downstream lesson. It is worth running before every commit.

Lesson 00 — Blinky

CONCEPT	A counter is a frequency divider.
YOU BUILD	The 8 LEDs counting up as a binary odometer.
NEEDS	Nothing. This is the start.

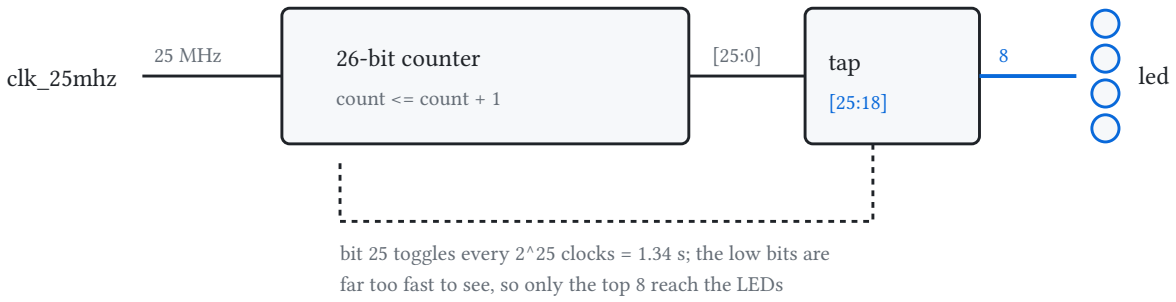
LEARNING GOALS

- Explain why bit n of a counter toggles every 2^n clocks.
- Use non-blocking assignment in a clocked block, and say why.
- Make a design's time constant a parameter, so it can be simulated.
- Read a resource report and account for what is in it.

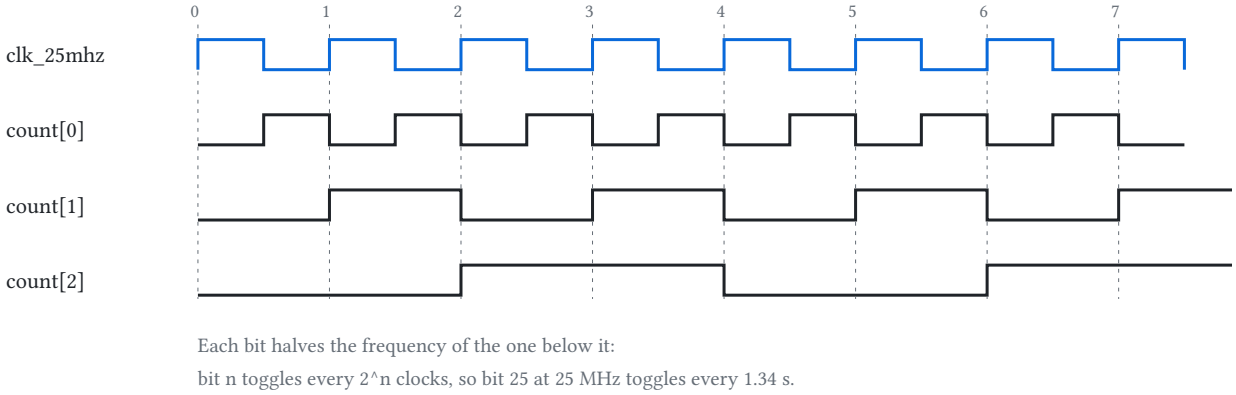
The idea

The board's crystal runs at 25 MHz. An LED toggling at that rate looks like an LED that is simply on, so to see anything the clock must be divided by tens of millions.

The cheapest divider in digital design is a binary counter. Add one to a register every clock and bit n toggles every 2^n clocks: bit 0 at 12.5 MHz, bit 1 at 6.25 MHz, and so on. Nothing is **built** to divide — the division is what counting already is, read one bit at a time.



Tapping the top eight bits of a 26-bit counter. At 25 MHz, bit 25 toggles every 1.34 seconds.



Each bit halves the frequency of the one below it — the whole of clock division, in one picture.

Theory

Non-blocking assignment

```
always @(posedge clk)
    count <= count + 1'b1;
```

⩽ is non-blocking: every right-hand side in the block is evaluated before any left-hand side updates. That is what a bank of real flip-flops does — all of them sample their inputs at the same instant. Lesson 03 measures what happens when you use = instead, and the answer is two flip-flops fewer than you asked for.

Parameterised time constants

The testbench overrides `WIDTH` from 26 to 8, so the simulation finishes in milliseconds rather than simulating 67 million clock cycles. **Making a design's time constant adjustable is what makes slow hardware testable**, and it is a habit every later lesson depends on — lesson 03's debouncer and lesson 09's baud generator both do it.

Part-select

```
assign led = count[WIDTH-1 -: 8];
```

"Eight bits, counting downward from `WIDTH-1`." Writing `count[25:18]` means the same thing until someone changes `WIDTH`, at which point it silently reads out of range.

What it cost

MEASURED a 26-bit counter

`TRELLIS_FF`	26	one flip-flop per counter bit, as expected
`CCU2C`	13	the ECP5's carry cells — **two bits each**
`LUT4`	1	

Thirteen carry cells do 26 bits of addition in dedicated hardware, which is why a wide counter is nearly free. The general-purpose logic contributes a single LUT.

WHY THIS MATTERS

This is the first appearance of the course's central method. The number is not "about thirty cells"; it is 26 flip-flops and 13 carry cells, and both are explainable. When a resource report contains something you cannot explain, you have inferred something you did not intend.

Learning outcomes

By the end you should be able to say, without looking anything up, why bit 25 of a counter clocked at 25 MHz toggles roughly every 1.3 seconds; why ⩽ belongs in a clocked block; and why a 26-bit counter costs one LUT rather than twenty-six.

Lesson 01 — Logic gates

CONCEPT	AND, OR, NOT, NAND — and the fact that this FPGA contains none of them.
YOU BUILD	Every two-input gate, verified against its complete truth table, plus NAND-only versions of four of them.
NEEDS	Lesson 00.

LEARNING GOALS

- Write combinational logic as continuous assignment, not as statements.
- Test exhaustively while the input space still allows it.
- Explain NAND universality, and why it mattered more in 1975 than now.
- Measure why a 2-input AND and a 4-input XOR cost the same.

The idea

A gate is a Boolean function of its inputs. In Verilog you write it as a permanent connection rather than as something that happens:

```
assign y_and = a & b;
```

`assign` is not a statement that runs. There is no "when". It is a wire: if `a` or `b` changes, `y_and` follows as fast as the silicon allows.

Theory

NAND is universal

Every Boolean function can be built from NAND alone:

function	from NAND	why
`NOT a`	`NAND(a, a)`	both inputs equal, so the AND is just `a`
`a AND b`	`NAND(NAND(a,b), NAND(a,b))`	the second NAND is a NOT
`a OR b`	`NAND(NOT a, NOT b)`	De Morgan: `a|b == ~(~a & ~b)`
`a XOR b`	four NANDs	the classic construction

Historically this mattered enormously: one gate type meant one mask, one part number, one bin of 7400s on the shelf.

Exhaustive testing

A two-input gate has four possible inputs. There is no need to **choose** test cases — check all of them, against truth tables written out by hand rather than computed with the same operator being tested.

> **WHY THIS MATTERS**
> This is the last lesson where exhaustive testing is this trivial, and lesson 02 is the last where it is possible at all. The counts go 4, then 512, then — by lesson 10 — 4.6×10^{19}. Build the habit while the habit is cheap.

What it cost, and why it is the point

> **MEASURED eleven outputs, seven LUT4s, zero flip-flops**
> Fewer LUTs than outputs, because the synthesiser does not implement what you **wrote**. It works out what each output *is*, as a function of `a` and `b`, and shares whatever can be shared.
>
> The NAND-only constructions are the clearest case. `n_xor` was written as a chain of four NANDs; the synthesiser evaluated the function, found XOR, and programmed one lookup table. **It never built the chain.**

The fabric has no gates — only 4-input lookup tables (page 7). So a 2-input AND and a 4-input XOR cost exactly the same: one LUT4 each. So does any other function of up to four inputs, however complicated.

WHY THIS MATTERS
Minimising gate count is half of a classic digital design course, and on an FPGA it is almost pointless. What costs is the number of LUTs, their *depth* — how many in series, which sets your clock speed — and the routing between them. Rewriting an expression to use three gates instead of four usually changes nothing at all.

The boundary that does matter

A LUT4 holds any function of four inputs. Five inputs needs at least two LUTs, no matter how simple the function:

function	cost
`(a & b) | (c ^ d)`	1 LUT4
`((a^b) & (c|~d)) | (~a&b&c&d) | (a&~b&~c)` — far uglier, same 4 inputs	1 LUT4
`(a & b) | (c ^ d) | (e & ~a)` — one more input	2 LUT4 + 1 mux

Count inputs, not operators. "Simplifying" a four-input expression from five operators to three saves nothing.

Learning outcomes

Why `assign` is a connection rather than a statement; why NAND is universal; why a 2-input AND and a 4-input XOR cost the same here; and the difference between `&` and `&&`, which is invisible on one-bit signals and produces wrong answers on a bus.

Lesson 02 — Combinational logic

CONCEPT	Muxes, decoders and adders — the three structures everything else is made of.
YOU BUILD	A 7-segment hex decoder and a parameterised adder, both verified exhaustively.
NEEDS	Lesson 01.

LEARNING GOALS

- Recognise the mux and the decoder as mirror images, and find them in a CPU.
- Explain why addition has a critical path at all.
- Avoid the two silent width traps: shift results and discarded carries.
- Measure why writing `+` beats building an adder from gates by 12×.

The idea

Real designs are not built from gates any more than prose is built from letters. Three structures account for most of a processor.

The multiplexer selects: one of several inputs reaches the output. Every point in a datapath where a value "comes from here, or from there" is a mux, and lesson 13's CPU is mostly muxes with arithmetic between them.

The decoder is the mux's mirror image: one input number, and exactly one of 2^n outputs goes high. Wherever a number must become a *choice* — instruction decode, address decode, one-hot state — there is a decoder.

The adder is the first structure with an interesting internal shape, because each bit's result depends on a carry from the bit below.

Theory

Carry propagation is why addition takes time

One bit of addition is two gates' worth of logic:

```
assign s    = a ^ b ^ cin;                // odd number of ones
assign cout = (a & b) | (cin & (a ^ b));  // at least two ones
```

Chain W of them and the top bit of the sum cannot settle until the carry has propagated through every stage below it. At 4 bits nobody cares. At 32 bits it is the slowest thing in the CPU — and lesson 10 measures exactly that.

Two width traps

Both are silent, and both come from the same root: **in Verilog, width is part of the meaning.**

> **WATCH OUT**
>
> **Shifts.** The result of `a << b` is as wide as `a`. So `assign y = 1'b1 << sel;` shifts a **one-bit** value, everything falls off the end, and `y` is always zero. The literal must be widened before it is shifted.
>
> **Carries.** `a + b + cin` assigned to a W-bit destination throws the carry away. Widening the intermediate to W+1 bits is what keeps it. At 4 bits this is wrong in exactly half of the 512 possible cases — and the sum bits are never corrupted, only the carry, which is what makes it hard to spot.

Latches

A combinational `always` block that fails to assign its output on some path infers a **latch** — a storage element with no clock, which is almost never what you meant. The habit is to assign a default first and then override:

```
always @(*) begin
    y = 1'b0;            // default FIRST, so no path is uncovered
    case (sel) ... endcase
end
```

What it cost

MEASURED a case statement and a ROM array are the same thing

`seg7` — a 16-way `case` statement	7 LUT4
`seg7_rom` — the same table as an array of constants	7 LUT4

The synthesiser cannot tell the difference, because there is none. A `case` is not "logic" and an array is not "memory"; both describe one function of four inputs, and only the function survives.

MEASURED hand-built addition versus the + operator

	at 4 bits	at 32 bits — CPU width
built from full adders	24 LUT4 + 17 muxes	138 LUT4 + 74 muxes
`a + b + cin`	**3 CCU2C**	**17 CCU2C**

At CPU width that is **212 cells against 17** — a factor of twelve.

The reason is what is *absent* from the hand-built rows: there are no `CCU2C` cells at all. The carry chain is dedicated hardware, one cell per two bits, and when you write `a + b` the synthesiser recognises addition and maps straight onto it. Write the same function as explicit gates and it sees gate soup — it cannot prove that pile of XORs and ANDs is an adder, so the hard carry logic goes unused.

WHY THIS MATTERS

Lessons 01 and 02 point in opposite directions and both are true. In lesson 01 the synthesiser *discarded* your structure and kept only the function. Here it *kept* your structure, because your structure prevented it from recognising the function.

The rule underneath both: **write what you mean, at the highest level you can, and let the tool match it to the hardware.** Build the ripple adder once to understand what a carry is. Then write `+` forever.

Learning outcomes

Why a mux and a decoder are mirror images; why `1'b1 << sel` is always zero; why a `case` statement and a lookup table cost the same; and why the hand-built adder loses by twelve times in terms of what the synthesiser can and cannot recognise.

Lesson 03 — Sequential logic

CONCEPT	The flip-flop — the only component that makes time exist.
YOU BUILD	A debounced button that increments a counter exactly once per press.
NEEDS	Lesson 02. A board makes the point far better.

LEARNING GOALS

- Explain `⩽` versus `=` as two different circuits, not as a style rule.
- Design a debouncer, and choose its time constant from first principles.
- Turn a level into a one-clock pulse, and know when you need to.
- Distinguish the missing `else` that is a latch from the one that is not.

The idea

Everything so far had no memory. A D flip-flop changes that: on each rising clock edge it samples its input and holds that value until the next edge.

Nothing else in this course stores a bit. A register, a shift register, a counter, a state machine and a 32-bit CPU are all arrangements of flip-flops with combinational logic deciding what goes into them.

Theory

`⩽` versus `=`, as two circuits

The same three-stage shift register, written twice:

```
always @(posedge clk) begin
    a ⩽ d;
    b ⩽ a;    // the OLD a
    c ⩽ b;    // the OLD b
end
```

```
always @(posedge clk) begin
    a = d;
    b = a;    // the NEW a
    c = b;    // the NEW b
end
```

Non-blocking evaluates every right-hand side before updating any left-hand side, so each stage samples what its neighbour held at the start of the edge. A pulse walks along: stage `a` after one clock, `b` after two, `c` after three.

Blocking takes effect immediately, so each line sees the result of the line above. All three registers load the same value on the same edge. The delay line is gone.

MEASURED the bug, as a number

`shift_nonblocking`	**3** flip-flops
`shift_blocking`	**1** flip-flop

One flip-flop. The synthesiser worked out that `b` and `c` always equal `a`, built one register and wired its output three ways.

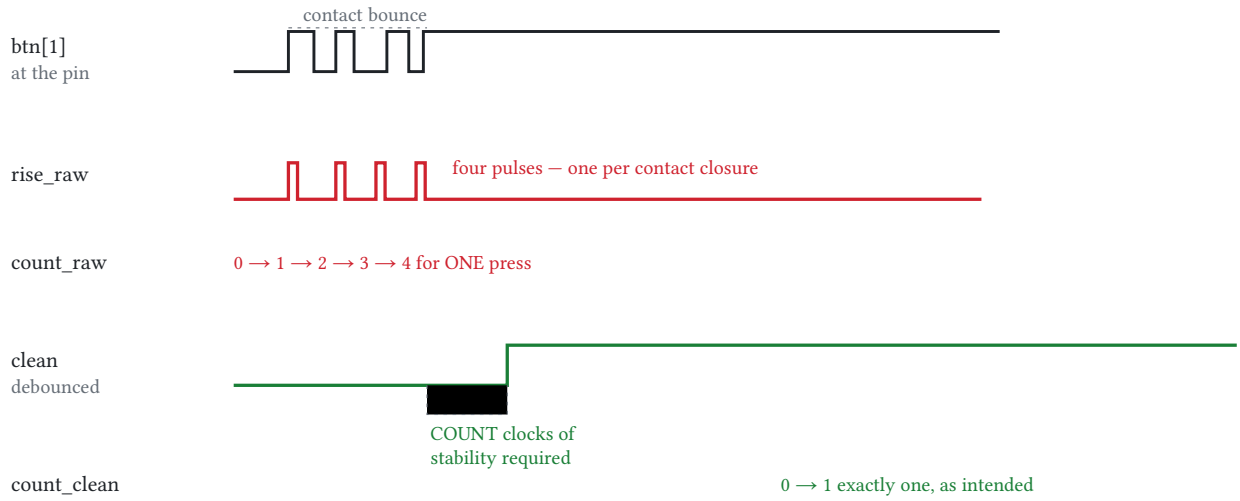
WATCH OUT

Nothing warns. It compiles, it simulates, it synthesises, and `make lint` is quiet. The three-stage delay you asked for became a one-stage delay, and the only ways to notice are to know the rule or to count the flip-flops. This is the most common serious mistake in Verilog.

A switch is not a digital input

A button is two pieces of springy metal. When they meet they bounce apart and meet again — several times, over roughly 1–10 ms.

At 25 MHz, 10 ms is **250,000 clock cycles**. The bounce is not a blur or a glitch to the FPGA: it is tens of thousands of clocks of perfectly clean, fully sampled, entirely real transitions. **One press genuinely is several presses**, as far as the hardware can tell.



One press. The raw path counts three times; the debounced path counts once.

Lesson 02 did not care — combinational gates just flicker, and a flickering LED is not a wrong answer. The moment an edge is used to *count*, every bounce counts.

The defence is a saturating counter: the input must disagree with the current output *continuously* for `COUNT` clocks before the output follows. Any shorter disagreement resets the timer, which rejects bounce of any pattern rather than only bounce shorter than one sample period. The price is latency, so you pick the smallest `COUNT` that comfortably exceeds the switch's bounce time.

The apparent contradiction

```
module reg_w #(parameter W = 32)(input clk, en, ...);
    always @(posedge clk)
        if (en) q ⩽ d;       // no else
endmodule
```

In lesson 02 an `always @(*)` block without an assignment on every path inferred a latch, and that was a bug. Here the identical-looking omission is **correct**.

The difference is the clock. A flip-flop already holds its value between edges, so "assign nothing" is a real, cheap, well-defined behaviour — it is just the flop's enable input. In a combinational block there is no storage to fall back on, so the tool has to invent some.

> **WHY THIS MATTERS**
> Same syntax, opposite meaning, decided entirely by whether the block is sensitive to a clock edge or to its inputs. This is the one place the two rules appear to conflict, and getting it straight here saves confusion in every later lesson.

What it cost

MEASURED every flip-flop accounted for

debounce: an 18-bit counter plus `clean`	19
two edge detectors, one `prev` each	2
two 4-bit counters	8
total	**29**

And the nine carry cells are the debounce counter's 18-bit increment — one per two bits, exactly as lesson 02 predicted.

Learning outcomes

Why `<=` gives three flip-flops and `=` gives one; why 10 ms of bounce is a quarter of a million real clock cycles; why a saturating counter beats a simple delay; and why `if (en) q <= d;` with no `else` is correct in a clocked block and a bug in a combinational one.

Lesson 04 — Clocks and timing

CONCEPT	What a clock actually is inside an FPGA, and what the tools check when they check timing.
YOU BUILD	A PLL generating 100 and 40 MHz, a safe crossing between two clock domains, and a design pushed until it fails timing.
NEEDS	Lesson 03. A board, for the PLL — it cannot be simulated at all.

LEARNING GOALS

- Explain why a design is correct *at a frequency* rather than correct.
- Divide a clock to make an *enable*, never a clock, and say why.
- Explain why a PLL's output frequencies are not freely chosen.
- Cross a bus between clock domains without corrupting it.
- Explain why no simulation can validate a clock-domain crossing.

The idea

Three lessons used `clk_25mhz` as if it were simply given. It is not.

Setup, hold, and the critical path

A flip-flop needs its data input stable *before* the clock edge (setup) and *after* it (hold). Between any two flip-flops there is a budget: one clock period, minus the first flop's output delay, minus the second's setup time. Everything else is available for logic and routing.

The longest such path in the design sets the maximum frequency:

$$\text{max frequency} = \frac{1}{\text{critical path delay}}$$

> **WHY THIS MATTERS**
> So **a design is not correct or incorrect in isolation — it is correct at a frequency.** The same Verilog that works at 25 MHz may fail at 100, and nothing in the source says which. This is also the real argument for pipelining, twelve lessons early: lesson 16 splits the CPU into five stages not for elegance but because a shorter critical path permits a faster clock.

Enables, not gated clocks

Lesson 00 divided the clock with a counter and used the result as *data*. Using it as a *clock* is the mistake worth naming:

```
wire tick = (div == LIMIT);                          // yes: an enable
always @(posedge clk_25mhz) if (tick) count <= count + 1'b1;

always @(posedge div[15]) count <= count + 1'b1;    // no: a derived clock
```

An FPGA has a small number of **global clock networks** — dedicated low-skew routing that reaches every flip-flop at very nearly the same instant. A signal generated in logic travels on ordinary routing, arrives at different flops at measurably different times, and the timing analyser can no longer give a meaningful answer. You can watch the allocation happen in the build log:

```
Info:  promoting clock net clk_100mhz to global network
Info:  routing clock net $glbnet$clk_100mhz using global 1
```

The PLL is a hard block

The ECP5's `EHXPLLL` is a real analogue phase-locked loop in the corner of the die, not logic. Its parameters are computed by a calculator rather than chosen:

```
ecppll -i 25 -o 100 --clkout1 40
  → VCO 600 MHz, divisors 6 and 15
```

The VCO must land in its legal range (roughly 400–800 MHz) and **every output is that one VCO frequency divided by an integer**. You cannot ask for an arbitrary pair of frequencies.

Three consequences:

1. **It cannot be simulated.** Icarus has no model and Verilator has no definition. This is the first lesson where what you simulate and what you synthesise differ — the lesson's Makefile names a simulatable subset, and the PLL is verified structurally and on hardware instead.
2. **`locked` matters.** The loop takes tens of microseconds to settle. Logic in a PLL domain must be held in reset until then.
3. **PLLs are finite.** Four on this part, not one per module.

Crossing between domains

Two clocks not derived from one another have **no fixed phase relationship**. Sooner or later a signal from domain A changes at exactly the wrong moment relative to domain B's edge — inside the setup/hold window. The receiving flip-flop then goes **metastable**: its output sits between 0 and 1 for an unbounded time before resolving arbitrarily.

> **WATCH OUT**
> **This cannot be simulated.** Verilog has no notion of setup time, so the simulator always gives a clean, defined, repeatable answer no matter how badly the crossing is built. Simulation cannot show you this bug and cannot prove you have fixed it.
>
> The defence is *structural*: you follow the rule because the rule is correct, not because a test went green. It is also why "it worked on my bench for an hour" means nothing — a mean time between failures of three weeks looks exactly like a working design on a Tuesday afternoon.

For **one bit**, a two-flop synchroniser. The first flop may go metastable; the second samples it a full period later, by which time it has almost certainly resolved. "Almost certainly" is the honest phrasing.

For **a bus**, a synchroniser per bit is a *bug*. Each bit resolves independently, so some arrive one clock and some the next, and the destination briefly sees a value that was never sent — `0x0F → 0x10` observed as `0x1F` or `0x00`. Instead, synchronise *one* bit and use it to say when the others are safe to read; the data itself crosses with no synchroniser at all, because by the time the request's toggle is noticed the data has been stable for two destination clock periods.

What was measured

> **MEASURED timing on the real build**

`clk_25mhz` domain	achieves 342.82 MHz — PASS at 25
`clk_100mhz` domain	achieves 278.09 MHz — PASS at 100
cross-domain path	1.49 ns, **found and reported separately**

nextpnr finds the crossing unaided, because it cannot analyse it as a normal path. An unexpected `→ posedge <other clock>` line in your report is the tool handing you a bug report.

> **MEASURED why 16× oversampling is not an option at 25 MHz**
> Lesson 09's UART wants to sample at 16 times the baud rate. But 115,200 × 16 = 1,843,200 Hz, and 25,000,000 / 1,843,200 = 13.5634 — not an integer. The nearest divisors give +4.33% and −3.12%, which is 41% and 30% of a bit adrift by the stop bit. The textbook technique is simply unavailable on this clock.

The honest part about forcing a failure

Two attempts to build something deliberately slow were defeated by the synthesiser. First a wrapper read only the low byte of the result, so the upper 24 bits were pruned and the measured frequency went *up* with chain

depth. Then, with every bit live, a 16-deep chain still collapsed to 8 carry cells and 9 LUTs — the arithmetic was folded.

WHY THIS MATTERS
You cannot read your critical path off your source code. The thing that gets built is not the thing you wrote. Lessons 01 and 02 showed that for area; it is equally true for delay. Read the tool's report.

The reliable way to see a failure is to ask for more speed, not to build something slower:

MEASURED the same placed design, asked for different clocks

asked for	achieved	verdict
100 MHz	357.65 MHz	PASS
300 MHz	357.65 MHz	PASS
400 MHz	357.65 MHz	**FAIL**

The achieved number never changes — it is a property of the placed and routed design. What changes is whether it is enough. That is what timing closure *is*.

Lesson 05 — Finite state machines

CONCEPT	Encoding "where have I got to" as a register, with transitions as combinational logic.
YOU BUILD	A traffic light, two sequence detectors contrasting Moore with Mealy, and a combination lock.
NEEDS	Lesson 04.

LEARNING GOALS

- Write an FSM in the canonical three-part form, and say why it is three parts.
- Choose between Moore and Mealy on the basis of what the consumer needs.
- Measure state encoding rather than inheriting a rule about it.
- Handle a wrong symbol in a sequence detector without losing a valid prefix.
- Test by transition coverage, since there is no input space to sweep.

The idea

A state machine is a register holding *where you are*, and combinational logic deciding *where you go next*. That is all, and everything structural follows from it:

1. a **state register** — one clocked block, containing nothing else
2. **next-state logic** — one combinational block, a pure function
3. **output logic** — from state (Moore), or state and inputs (Mealy)

Splitting (1) from (2) is not stylistic. It makes the only sequential element trivially obvious, so the next-state block reads as a plain truth table and you never wonder what is registered.

> **WHY THIS MATTERS**
> The SPI engine in lesson 06, the UART in lesson 09 and the CPU's control in lesson 13 are all this one pattern at increasing scale. The CPU is not a different kind of thing; it is a state machine whose state includes a program counter.

Draw the diagram first, on paper, before any Verilog: bubbles for states, arrows labelled with the condition that takes them. Skipping that step is how you end up with a state you cannot reach and a transition you cannot explain.

Moore and Mealy

Two machines detecting the same sequence, differing only in where the output comes from:

	Moore	Mealy
output depends on	state	state **and** input
states needed	4	3
timing	one clock **late**	**immediate**
output is	registered, glitch-free	combinational

Neither is correct in general. The question is whether the consumer needs the answer *now* or needs it *clean*. A Mealy output feeding another FSM's next-state logic in the same cycle is a combinational path between two state machines — occasionally what you want, frequently a timing problem.

Two textbook rules, measured

MEASURED **Mealy has fewer states and costs more**

`seq_moore`	2 LUT4, 2 flip-flops
`seq_mealy`	**5 LUT4**, 2 flip-flops

Mealy genuinely needs fewer states — 3 against 4 — and the saving buys nothing, because both counts fit in two flip-flops, while making the output depend on the input adds real gates. The textbook trade is an *ASIC* trade, where you count flops.

MEASURED one-hot encoding, which is supposed to suit FPGAs, loses badly

encoding	cost
the synthesiser's own choice	11 LUT4, 3 muxes, 5 FF
`"binary"`	12 LUT4, 3 muxes, 5 FF
`"one-hot"`	**41 LUT4**, 30 muxes, 7 FF

Nearly four times worse. Four states is too few for one-hot's decode saving to appear, while every comparison against the state must now be built from four bits instead of two. One-hot earns its keep on *wide* state machines, not narrow ones.

In Verilog the encoding is an attribute on the state register — not on the module, where it does nothing. Without it the synthesiser re-encodes as it sees fit, so the constants you wrote are a suggestion rather than an instruction.

WHY THIS MATTERS
That is twice in two lessons that a received rule did not survive measurement. The method is the content: `make synth` takes two seconds and outranks anything anyone tells you about what *usually* costs less.

The mistake everyone makes

A lock opening on the sequence **1, 2, 1, 3**. The states are easy. The interesting question is *what happens on a wrong symbol.*

The obvious answer — go back to the start — is wrong. Enter `1, 2, 1, 2`: after `1, 2, 1` the machine wants `3` and gets `2`. But the last two symbols it saw were `1, 2`, which is a valid *prefix* of the code. Restarting throws that away, and the lock then refuses `1, 2, 1, 2, 1, 3` even though the correct code is right there at the end.

Every state needs the same treatment: on a wrong symbol, fall back to the longest suffix of what you have seen that is still a prefix of the code. Doing it by hand is fiddly and easy to get subtly wrong — it is precisely what the Knuth–Morris–Pratt string-matching algorithm automates, and precisely the class of bug lesson 14's formal tools find without being asked.

Testing changes shape

An FSM's input is a *sequence*, so there is no space to sweep. The test is a set of chosen journeys through the state graph, and the discipline is **transition coverage**: has every arrow in your diagram been taken at least once?

The traffic-light test walks the full cycle for that reason, and then runs many clocks with no ticks to confirm the machine does not advance without its enable.

What it cost

MEASURED 113 flip-flops, and the proportion is the lesson

one-second tick divider (25 bits)	25
traffic light: 2 state + 3 timer	5
four debouncers	76
four edge detectors	4
lock state	3
total	**113**

The state machines are 8 flip-flops out of 113. The rest is debouncing and timekeeping — the plumbing needed to talk to the physical world. That ratio is normal, and worth seeing once: the interesting logic is usually a small fraction of a real design.

Lesson 06 — SPI and the OLED

CONCEPT	A state machine that talks to a real chip, on that chip's terms.
YOU BUILD	The SSD1331 colour OLED brought up from cold and filled with an animated gradient — with no framebuffer anywhere.
NEEDS	Lesson 05. A board, to see it.

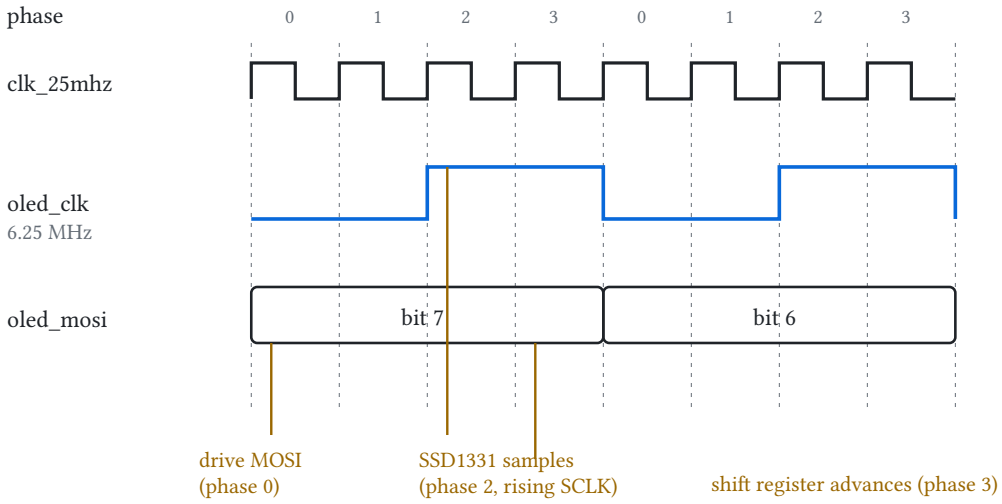
LEARNING GOALS

- Read a datasheet's timing requirements and turn them into clock counts.
- Set up data before clocking it, and say why that ordering is universal.
- Use a `busy`/`start` handshake to separate timing from sequencing.
- Recognise when a design needs no frame buffer at all.

The idea

Everything so far was self-contained: your logic, your clock, your decisions. This is the first lesson where another piece of silicon has opinions, written down in a datasheet, and your design must satisfy them.

The SSD1331 speaks 4-wire SPI. The contract is: you supply the clock, at no more than about 6.6 MHz; data goes out most significant bit first; `D/C` is a *level* rather than a bit in the stream, saying whether this byte is a command or pixel data; and `RES#` must be held low on power-up and then released with a real delay before anything else.



Four system clocks per bit: data is set up two clocks before the chip samples it.

One byte. Four system clocks per bit gives 25 MHz ÷ 4 = 6.25 MHz, inside the datasheet's limit.

Theory

Set up the data, then clock it

The four phases do distinct jobs: drive the data line on phase 0, raise the clock on phase 2 so the data has been stable for two clocks before the chip samples it, and advance the shift register on phase 3.

> **WHY THIS MATTERS**
> "Set up the data, *then* clock it" is the whole of synchronous interface design. The same reasoning sets the sample point in lesson 09's UART receiver, where the receiver cannot control the clock and must infer the safest moment to look. Phase 1 in this design does nothing at all, and that idleness is the setup time.

The handshake

```
if (!busy && !start) begin ... end
```

`busy` is the entire interface between the bit-level engine and everything above it. The sequencer that pushes initialisation bytes and pixel data obeys one rule — do not touch `data` or raise `start` while `busy` is high — and therefore needs to know *nothing* about SPI timing.

That separation is why the same sequencer survives unchanged into lesson 07, which replaces only the pixel function.

> **WATCH OUT**
> `start <= 0;` sits at the top of the sequencer's always block, before the case. That makes `start` a one-clock pulse by default which any branch may override. Forget it and `start` stays high, the engine re-triggers immediately, and you send one byte twenty-five million times a second.

No framebuffer

96 × 64 × 16 bits is 98,304 bits, which would fit in block RAM easily. It is not needed, because the SSD1331's internal address pointer auto-increments: set the column and row window once during initialisation, and every subsequent data byte lands in the next pixel automatically.

So the design only has to produce bytes in raster order, and `x`/`y` counters mirroring the chip's own pointer are enough. The pixel is *computed* from the coordinate the moment it is asked for:

```
wire [15:0] pixel = { x[6:2], y[5:0], frame[7:3] };
```

> **WHY THIS MATTERS**
> Recognising that a buffer is unnecessary is worth as much as knowing how to build one. Lesson 08 builds the framebuffer that lesson 15's software will need — but a design that can compute its output on demand should.

What it cost

MEASURED the whole display controller

LUT4	108
TRELLIS_FF	72
CCU2C	22
PFUMX	29

The 72 flip-flops are the shift register, the phase and bit counters, the state, the reset delay counter, and the x/y/frame counters. The 29 multiplexers are mostly the 31-entry initialisation ROM, which the synthesiser built out of logic because it is far too small to be worth a block RAM — a decision lesson 08 examines properly.

Testing an interface

The testbench decodes the SPI stream the way the OLED would: it samples the data line on each rising clock edge, assembles bytes, checks the first 31 against the expected initialisation sequence with `D/C` low, then confirms pixel data keeps arriving with `D/C` high.

> **WHY THIS MATTERS**
> It reconstructs bytes *from the wire* rather than peeking at internal signals, so it tests the contract the other chip cares about. A bit-order or clock-phase error that an internal check would miss fails here.
>
> But note the limit: a passing simulation does not prove the design meets a datasheet. Halve the clock divider and the simulation still passes while the hardware probably will not. When another chip's timing is involved, simulation is necessary and not sufficient.

Lesson 07 — Font ROM and coordinate maths

CONCEPT	A lookup table plus arithmetic on coordinates, instead of stored pixels.
YOU BUILD	"HELLO" in white, scaled 2×, over the lesson 06 gradient — still with no framebuffer.
NEEDS	Lesson 06.

LEARNING GOALS

- Turn a raster coordinate into a ROM address and a bit select.
- Choose constants so that the arithmetic disappears.
- Recognise a hidden coupling between two distant constants.
- Explain why a 32-byte ROM becomes logic rather than memory.

The idea

Lesson 06 established that you can drive a display without storing a frame, as long as you can *compute* the pixel at any coordinate the moment it is needed. Text is the first interesting thing to compute.

A bitmap font is a tiny ROM: each character is eight bytes, each byte a row, each bit a pixel. So "is this pixel part of a letter?" becomes a chain of arithmetic on the current raster coordinate — which character cell is x inside, which glyph is in that cell, which row does y select, which bit does x select.

WHY THIS MATTERS
That chain is exactly what a hardware text-mode display controller does, and what a GPU's texture unit does in more dimensions: a coordinate in, an address and a bit-select out.

Theory: make the arithmetic disappear

```
wire [2:0] ci    = tx[6:4];        // which character cell: divide by 16
wire [2:0] frow  = y[3:1] - 3'd4;  // which glyph row: divide by 2
wire [2:0] fcol  = tx[3:1];        // which column: divide by 2
wire font_bit = font[{glyph, frow}][3'd7 - fcol];
```

Every one of those is **free**, and that is the point:

- `tx[6:4]` is a divide by 16 — a wire rename, no logic at all. Choosing a 16-pixel cell width (a power of two) is what makes it free; 10 pixels would have cost a divider.
- `tx[3:1]` is the same trick for the 2× scale: dropping the low bit makes every source pixel cover two screen pixels.
- `font[{glyph, frow}]` concatenates a 2-bit glyph index and a 3-bit row into a 5-bit address. Address arithmetic by *concatenation* rather than multiplication, because the glyph stride is 8.
- `3'd7 - fcol` reverses the bit order, because bit 7 of the font byte is the *leftmost* pixel while x increases rightward.

WHY THIS MATTERS
"Choose your constants so the arithmetic disappears" is one of the highest-leverage habits in hardware design. The same design with a 10-pixel cell and a 3× scale is functionally identical and dramatically more expensive. Lesson 08 meets the same idea with a price attached, when a framebuffer's `y × 96 + x` turns out to cost two hardware multipliers.

WATCH OUT
`frow = y[3:1] - 3'd4` quietly relies on the text band starting at $y =$ 24. Change the band's position and this constant must change with it — and

the two facts are forty lines apart with nothing connecting them. A constant derived from another constant should be *derived*, not duplicated.

What it cost, and the question it raises

MEASURED adding text to lesson 06

	lesson 06	lesson 07
LUT4	108	**185**
TRELLIS_FF	72	72 — **unchanged**
PFUMX	29	50

No new flip-flops at all. The text added no state, only combinational arithmetic. The extra 77 LUTs are the coordinate chain and the font ROM.

And there is the puzzle this lesson ends on: **a 32-byte ROM built out of logic, not memory.** The part has 208 block RAMs sitting idle and the synthesiser used none of them.

The reason is size. One block RAM holds 18,432 bits; spending a whole one on 256 bits would waste 98.6% of it. The decision flips somewhere between 32 bytes and a few kilobytes — and knowing roughly where is a useful instinct, which is what lesson 08 is for.

On what the test does not check

The testbench verifies that the SPI stream is well-formed and that the initialisation sequence is correct. It does *not* check that the letters are the right shape — that would mean reconstructing a 96 × 64 image and comparing against a reference.

WHY THIS MATTERS

This gap is deliberate and worth sitting with. **A passing test tells you only what it was written to ask.** The course returns to this repeatedly: lesson 09 finds a testbench asserting something false, lesson 11 finds a mutation that survives for a good reason, lesson 12 finds a hole in its own vectors, and lesson 13 finds a bug that 161 passing checks had missed.

Lesson 08 — Memories

CONCEPT	You do not instantiate a block RAM — you write a shape the synthesiser recognises, and the tidier-looking shapes cannot be recognised.
YOU BUILD	The four memory configurations the fabric offers, a ROM loaded from a file, a dual-port frame-buffer, and a memory self-test that runs on the board.
NEEDS	Lesson 07.

LEARNING GOALS

- Write the coding pattern that infers block RAM, and spot the ones that cannot.
- Explain why a combinational read rules out block RAM at any size.
- Know roughly where the size threshold between logic and memory sits.
- Present address, write-enable and write-data in the same cycle.

The idea

The part has 208 `DP16KD` blocks — 18,432 bits each, about 3.7 Mbit in total, true dual-port, with a **registered output**. You cannot instantiate one in this course's style; you write an array and the synthesiser infers it, under three conditions:

- the array is a `reg [W-1:0] mem [0:DEPTH-1]`
- reads and writes happen inside `always @(posedge clk)`
- **the read result goes into a register**

That last one is not a style preference. A block RAM's output is registered *inside the block*, physically. A combinational read is something a `DP16KD` cannot do, so asking for one takes block RAM off the table — permanently, at any size.

The measurement

MEASURED the same array, two ways, at W = 32

depth	synchronous read	combinational read
16	8 DPR16X4 + 32 FF	**8 DPR16X4**
64	**1 DP16KD**	32 DPR16X4 + 132 LUT4 + 96 muxes
256	**1 DP16KD**	128 DPR16X4 + **919 LUT4** + 672 muxes
1024	**2 DP16KD**	—

Three readings, and the third is the one people miss.

At 256 entries, one line of difference costs 919 LUTs. The combinational version becomes *distributed* RAM — 16 words of 4 bits built from a logic cell's lookup tables — plus an enormous multiplexer tree.

At 1024 × 32 the synchronous version is two block RAMs and nothing else. 32,768 bits, and a `DP16KD` holds 18,432. Clean inference, zero LUTs.

And at 16 entries both give distributed RAM, with the combinational one *cheaper* — it needs no output register. Small asynchronous arrays have a legitimate home, and that is exactly how a 16-entry register file gets built.

WHY THIS MATTERS

So the rule is not "always synchronous". It is **synchronous once it is big enough to want a block, and know where that line is.** Lesson 11 depends on this: RV32I's 32-entry register file needs a combinational read for a single-cycle CPU, and that turns out to be the right choice rather than a compromise.

Lesson 07′s puzzle, answered

MEASURED identical code, different sizes

`font_small`	32 bytes — lesson 07′s font	9 LUT4
`font_full`	1024 bytes	**1 DP16KD**

Lesson 07′s font became logic because it was *too small to be worth a block*. The right conclusion is not "use logic for small memories" — it is: **write the memory in the shape that can become a block RAM and let the synthesiser decide on size.** Writing a shape that cannot removes the decision for good.

`$readmemh` runs at elaboration time

```
initial $readmemh("font.hex", mem);
```

The synthesiser reads the file *during synthesis* and bakes the contents into the bitstream. The hex file is therefore a **source file** — change it and you must resynthesise. This is the path lesson 13 uses to load a program, and lesson 15 uses for a program compiled by GCC.

WATCH OUT
A ROM whose contents are all zero is not a ROM — it is a constant, and the synthesiser folds it away entirely. The lesson's font file begins with ASCII 0–31, which have no glyphs, so synthesising it at 256 entries emits *no memory at all*. "My ROM disappeared" usually means the file was not found, was empty, or was genuinely all one value across the range used.

Two traps found while building it

A BUG THIS COURSE ACTUALLY HIT
Registered write data. The memory self-test registered `din` inside the clocked block, so by the time the RAM performed the write on the following edge, `addr` had already advanced. *Every location held its neighbour's value* and the read-back reported 128 errors out of 128.

A memory's address, write-enable and write-data must be presented in the **same cycle**. Register some and not others and they arrive skewed — and the memory faithfully stores the skew.

A BUG THIS COURSE ACTUALLY HIT
A constant multiply that was not free. The framebuffer's address is `y × 96 + x`. A comment in the first draft predicted the synthesiser would reduce a constant multiply to shifts. It did not — it allocated **two `MULT18X18D` hardware multipliers**. Writing `(y << 6) + (y << 5)` explicitly, since 96 = 64 + 32, removes them entirely.

`wy * 96 + wx`	1 DP16KD + **2 `MULT18X18D`** + 8 CCU2C
`(wy<<6) + (wy<<5) + wx`	1 DP16KD + 14 CCU2C + 24 LUT4

Third time in this course that the synthesiser did not do what the received rule said. If you know the cheap form, write the cheap form — then check.

WATCH OUT
Uninitialised memory reads as **X** in simulation. The ECP5 happens to zero its block RAM from the bitstream, but Verilog does not promise it, so a test that relies on initialisation is testing the simulator. The lesson's framebuffer test writes its zeros explicitly, after an earlier version failed for exactly this reason.

Why a framebuffer wants two ports

A display memory is written one way and read another, simultaneously: written by coordinate in whatever order the drawing code likes, and read in raster order relentlessly, because the display never stops. One port cannot serve both without one side waiting, which is largely why the `DP16KD` exists at all.

Lesson 09 — UART

CONCEPT	Agreeing about time with a device you share no clock with.
YOU BUILD	A transmitter, a receiver, and an echo server you can type into from a terminal.
NEEDS	Lesson 05, and lesson 04 for the synchroniser. A board — this is where the FPGA talks to your computer.

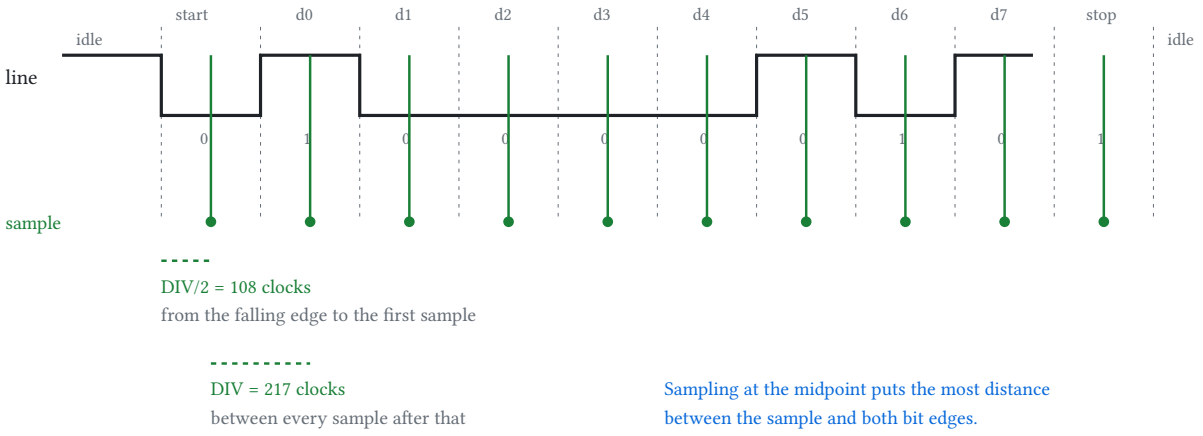
LEARNING GOALS

- Explain why a start bit exists when there is no clock on the wire.
- Compute a baud divisor and judge whether its error matters.
- Decide where in a bit to sample, and why.
- Explain why a framing error is not a reliable corruption detector.

The idea

There is **no clock on the wire**. Two chips, two crystals, one signal line each way. Everything depends on both ends agreeing in advance how long a bit lasts, and on one edge to say when to start counting.

That is the whole of 8N1: a start bit (the line falls from idle-high), eight data bits least significant first, and a stop bit. Ten bit periods, no handshake, no acknowledgement.



An 8N1 frame, with the receiver sampling at the midpoint of every bit.

Baud rate is a counter modulus

$$25,000,000 / 115,200 = 217.0139$$

So the divisor is 217, giving 115,207 baud — **0.0064% fast**. Over the 9.5 bit periods to the middle of the stop bit that accumulates to 0.06% of a bit, which is nothing. Even 216 or 218 would be fine, at 4.5% of a bit.

Where the textbook had to be overruled

Every UART reference says: build a clock at 16 times the baud rate and sample on its eighth tick. That is how a 16550 works, because a 16550 is *given* a 16× baud clock. Here it does not work.

MEASURED 16× oversampling is unavailable on a 25 MHz crystal

115,200 × 16 = 1,843,200 Hz, and 25,000,000 / 1,843,200 = 13.5634 — not an integer.

divisor	baud	error	drift by the stop bit
13	120,192	+4.33%	**+41% of a bit**
14	111,607	−3.12%	**−30% of a bit**

Either way the stop bit gets sampled inside bit 7 and the frame is lost.

WHY THIS MATTERS
The 16× clock is a good idea *given a crystal chosen for it*, and a bad idea given this one. So: count the bit period directly — which divides cleanly — and sample at the half-way point. Same underlying idea, which is to put the sample as far from both edges as possible, implemented to fit the clock you actually have.

The receiver re-checks the start bit at its own midpoint, which is what makes a glitch harmless: a brief disturbance looks like a start bit for a few clocks, and confirming it at the midpoint rejects it rather than reading eight bits of noise.

WATCH OUT
The receive line comes straight from a pin, driven by *another chip's oscillator*. It is asynchronous, so it goes through a two-flop synchroniser before anything looks at it. Lesson 04 exists for this moment, and this is the first lesson that genuinely needs it.

The honest limit

A receiver must absorb some rate error, because two crystals never agree. Driving frames at a range of rates against a correctly-configured receiver:

MEASURED tolerance, and what lies beyond it

sender error	result
+6.4% … −4.8%	accepted, byte correct
beyond that	**accepted, byte WRONG, no framing error**

The out-of-tolerance frames are *not rejected*. Far enough off, the stop-bit sample lands somewhere inside the data, finds a 1 by chance, and the frame passes its only check.

WHY THIS MATTERS
A framing error is not a reliable corruption detector. It catches many cases; it does not catch all of them. Detecting the rest is a *protocol* job — a parity bit, or a checksum over a packet. Hardware that cannot distinguish "correct" from "plausible" needs software that can.

A BUG THIS COURSE ACTUALLY HIT
The lesson's testbench originally asserted that mistimed frames were never silently corrupted — and passed. The sweep above disproves that claim. **A test that asserts something false and passes is worse than no test**, so it now asserts what is true and the tolerance window is documented with its numbers instead.

What a loopback test does not prove

Running all 256 byte values through transmitter into receiver is worth having — it finds bit-order and framing mistakes. But both ends share a clock and a divisor, so it **cannot find a timing problem**, which is the entire difficulty of a UART. The tolerance sweep is the test that does.

A BUG THIS COURSE ACTUALLY HIT
A second testbench bug, worth knowing because the pattern recurs. The bit-period measurement asserted `start` and *then* set up a wait for the line's falling edge. The edge had already happened, so it measured the next one — the five-zero run inside the letter `A` — and reported 40 clocks for an 8-clock bit. The transmitter was correct; the test was looking in the wrong place. Measuring edges in a testbench needs `fork`/`join` so the monitor is already waiting before the stimulus starts.

What it cost

MEASURED the asymmetry has a price

`uart_tx`	23 LUT4, 22 flip-flops
`uart_rx`	**53 LUT4**, 33 flip-flops

The receiver costs more than twice the transmitter. That is the resource footprint of not knowing when the bits start: the transmitter only has to count, while the receiver has to work out *when to look*.

WATCH OUT

The pin names are from the **FTDI's** point of view. `ftdi_rxd` is where the FTDI *receives*, so the FPGA drives it; `ftdi_txd` is where the FTDI transmits, so the FPGA reads it. `output ftdi_rxd` and `input ftdi_txd` are correct and look backwards. This is the classic place to lose an afternoon.

Lesson 10 — The ALU

CONCEPT	32-bit arithmetic and logic as one combinational block — and the point where exhaustive testing becomes impossible.
YOU BUILD	The complete RV32I ALU, verified against a reference model and an independent second implementation.
NEEDS	Lesson 08.

LEARNING GOALS

- Explain why RV32I has exactly ten integer operations.
- Derive subtraction, signed and unsigned comparison from one adder.
- Get signed less-than right, including where the difference overflows.
- Replace exhaustive testing with directed cases, randomisation and an independent model.
- Verify that a testbench has teeth, by mutating the design.

The idea

From here the modules are not illustrations. They are the processor's parts.

RV32I's integer ALU has **ten operations and not one more**: add, sub, three shifts, two comparisons, and three bitwise operations. A CISC machine might have hundreds. Everything else is built by the compiler out of these — negation is `sub` from `x0`, bitwise NOT is `xor` with −1, branch-if-greater is branch-if-less with the operands swapped.

That austerity is the "reduced" in RISC, and it is a claim about what *hardware* should cost rather than about what programmers should write.

The encoding is a gift

The operation selector is `{instruction[30], funct3}` — four bits. Only **two** operations need bit 30, and they are exactly the two pairs where one is the signed or arithmetic variant of the other: `ADD`/`SUB` and `SRL`/`SRA`. So lesson 12's decoder pulls this field straight out of the instruction word with no translation table at all.

Theory

One adder does four jobs

`a - b` $=$ `a + (~b) + 1` in two's complement, so subtraction is the *same* adder with `b` inverted and the carry-in set. And both comparisons fall out of that same subtraction for free:

Unsigned less-than is the borrow. Carry-out is 1 when there was no borrow, i.e. when `a` $\geqslant$ `b`. So `a < b` is simply `~carry` — no comparator at all.

Signed less-than is where ALUs go wrong. It is not a comparison of the raw bits, and it is not simply the sign of the difference either, because the difference can overflow:

```
wire overflow = (a[W-1] ^ b[W-1]) & (sum[W-1] ^ a[W-1]);
wire lt       = sum[W-1] ^ overflow;
```

> **WATCH OUT**
> The shortcut that nearly works: "if the signs differ, the answer is just `a`'s sign; a negative number is always less than a positive one." That is true, and implementing it that way passes almost every test you would write by hand. It fails exactly where the subtraction overflows — `INT_MIN` against `INT_MAX`, and `0` against `INT_MIN`.

The shift trap

```
wire [W-1:0] sra = $signed(a) >>> sh;
```

`>>>` is an arithmetic shift **only if its left operand is signed**. Writing `a >>> sh` on an unsigned wire silently performs a *logical* shift. It compiles, simulates, lints clean, and is wrong for every negative input. `$signed(a)` is the whole fix, and this is the most common ALU bug there is.

Exhaustive testing dies here

A 32-bit ALU has $2^{32} \times 2^{32} \times 10$ input combinations — about 4.6×10^{19}. At a billion a second, 1,460 years.

What replaces it is the pattern for everything from here on:

1. **Directed edge cases.** Twelve interesting operands — zero, one, −1, `INT_MAX`, `INT_MIN`, `INT_MIN+1`, −2, alternating patterns, shift amounts 31 and 32 — crossed against each other for all ten operations. 1,440 vectors covering every extreme against every other extreme.
2. **Randomised comparison against a reference model.** 20,000 vectors. The model uses Verilog's operators and `$signed`; the design builds the same answers from an adder and explicit overflow logic. Two different routes to one answer.
3. **An independent second implementation, where a model would prove nothing.** This is the subtle one. The design's shifts *are* Verilog's shift operators, so checking them against the same operators tests nothing whatsoever. They are cross-checked instead against an iterative shifter that moves one place per clock — a genuinely different implementation that can disagree.

Mutation testing

A test that has never caught a bug is unproven. So: break the design deliberately and check the test notices.

MEASURED three mutations, three catches, with the exact operands named

`sra = a >> sh` — drop `$signed`	caught: `a=ffffffff b=1 → 7fffffff`
`lt = sum[W-1]` — drop overflow	caught: `a=0 b=80000000 → 1`
`ltu = carry` — wrong polarity	caught: `a=0 b=1 → 0`

WHY THIS MATTERS

Mutation testing is the only way to know a testbench has teeth. Lesson 09 found a testbench asserting something false and passing; this is the systematic version of that check, and every testbench from here on is worth mutating once.

The surprise in the resource report

MEASURED the shifter is about 80% of the ALU

`adder` (32-bit)	**17 CCU2C** — essentially free
`shift_barrel`	667 LUT4 + 401 muxes
`shift_iterative`	76 LUT4 + 40 flip-flops
the whole ALU	**885 LUT4 + 408 muxes + 17 CCU2C**

The adder — what you would assume an ALU is mostly about — is seventeen carry cells, because the fabric has dedicated carry hardware. Shifting has *none*, so a barrel shifter is a tree of multiplexers five layers deep and 32 wide, three times over for the three shift kinds.

The iterative shifter is **seven times smaller** and takes up to 31 cycles. On an 8-bit microcontroller that is the usual choice. In a single-cycle CPU it is not available at all, so lesson 13 pays for the barrel shifter whether it likes it or not.

WHY THIS MATTERS

"Which operations are expensive" is a property of the fabric, not intuition. Addition is free here and shifting is not; on a chip with no carry chain the answer would invert.

Timing: the ceiling lesson 13 will hit

A registered 32-bit ALU closes at **115.94 MHz**, with the critical path running through the adder's carry chain *and* the ten-way result multiplexer. Both on one path — which is precisely the argument lesson 16 makes for splitting them across pipeline stages.

Lesson 11 — The register file

CONCEPT	Two reads and one write per cycle from one array, and the one register hardwired to zero.
YOU BUILD	RV32I's 32 × 32-bit register file, plus a block-RAM alternative, measured against each other at two very different sizes.
NEEDS	Lesson 10.

LEARNING GOALS

- Explain why RISC instructions have at most two source registers.
- Explain what `x0` buys, and count the instructions it eliminates.
- Justify a combinational read here, having been warned against one in lesson 08.
- Explain why read-during-write must return the old value.
- Distinguish a mutation that survives from a test that is weak.

The shape is dictated by the instruction set

```
add rd, rs1, rs2
```

One instruction, **two source operands read in the same cycle**, one result written. So: two read ports and one write port. If the file had three read ports, instructions could have three sources. If it had one, R-type instructions as we know them could not exist.

> **WHY THIS MATTERS**
> The instruction set and the register file are the same design decision seen from two sides. Register-file ports are expensive, and that cost is visible in the shape of every RISC instruction set — this is the clearest example in the course of hardware constraints writing software syntax.

`x0`

Reads of `x0` return zero; writes are discarded. Two lines of Verilog, and they simplify the whole instruction set. None of the following are real instructions — they are assembler conveniences that exist *because* one register reads as zero and swallows writes:

what you write	what it assembles to
`mv rd, rs`	`addi rd, rs, 0`
`li rd, imm`	`addi rd, x0, imm`
`nop`	`addi x0, x0, 0`
`not rd, rs`	`xori rd, rs, -1`
`neg rd, rs`	`sub rd, x0, rs`
`beqz rs, L`	`beq rs, x0, L`
`j L`	`jal x0, L` — throw the return address away
`ret`	`jalr x0, 0(x1)`

One register sacrificed; eight instructions avoided. `jal x0, L` is the neatest: a jump-and-link whose link register is the bit bucket *is* an unconditional jump, so the ISA needs no `j` opcode at all.

Combinational reads, and why lesson 08 agrees

A single-cycle CPU must read its operands, compute, and write back inside one clock period, so the read has to be combinational. Lesson 08 showed that rules out block RAM — which sounded like a warning, and here is the correct answer, because lesson 08's *other* finding applies: at a few tens of entries the fabric builds arrays from lookup tables as distributed RAM, cheaply and with no latency.

> **MEASURED** at 32 registers, asking for block RAM does not even get you block RAM

combinational read	266 LUT4 + 198 muxes + 32 DPR16X4
registered read	214 LUT4 + 108 FF + 32 DPR16X4 — **no DP16KD at all**

Thirty-two entries is too small for the synthesiser to spend a block on, so the "block RAM version" got distributed RAM too — the same storage, plus 108 flip-flops of pipeline registers, plus a cycle of latency, for nothing.

MEASURED at 1024 registers the trade flips decisively

combinational read	**3,660 LUT4** + 454 muxes + 1,024 DPR16X4
registered read	**4 DP16KD** + 192 LUT4 + 131 FF

Same two modules, same code, opposite conclusion — decided entirely by size. A vector machine's register file lives on the right of that table; RV32I's lives firmly on the left.

Read-during-write returns the old value

During cycle *N* an instruction reads its operands, computes, and its result is written by the clock edge at the *end* of cycle *N*. So a read of the register being written returns the **old** value — which is not a limitation but a requirement: an instruction must see its operands as they were before it ran.

WATCH OUT

This confuses people in both directions. Some expect a write-to-read bypass and are surprised by the old value; others add one and silently change what their CPU computes. The testbench asserts the old value explicitly, so adding a bypass later is a visible, deliberate change.

Bypass becomes *necessary* when a read and a write belong to different instructions in flight at once — lesson 16's pipeline, and the block-RAM variant above, where latency forces the issue.

A mutation that should survive

Four mutations; three caught. The fourth — *allowing writes to* `x0` — survives, and that is the interesting result.

It survives because `x0`'s behaviour is enforced in **two independent places** and either one alone suffices. With the write guard gone, `regs[0]` holds garbage that the read mask makes unobservable. Remove *both* and the test fails immediately.

WHY THIS MATTERS

So the correct reading is not "the test is weak" but "the code is deliberately redundant" — and no behavioural test can distinguish them, because there is no behavioural difference. A test reaching inside to inspect `regs[0]` would be testing the implementation rather than the contract.

A surviving mutation means either your test missed something, or your code has belt-and-braces redundancy. Check which before fixing anything. Lesson 12 meets the other kind.

Exhaustive testing is briefly back

A register file has 32 addresses, not 2^{32} values, so the thing that matters — *which register* — can be swept completely. All 961 pairs of distinct registers are read simultaneously, which is what catches a second port quietly sharing the first one's address decoder.

Lesson 12 — Decode and control

CONCEPT	Turning a 32-bit instruction word into control signals — and why the immediate encoding looks deranged and is not.
YOU BUILD	A decoder for all of RV32I plus the immediate generator, tested against 52 real instruction encodings.
NEEDS	Lesson 11.

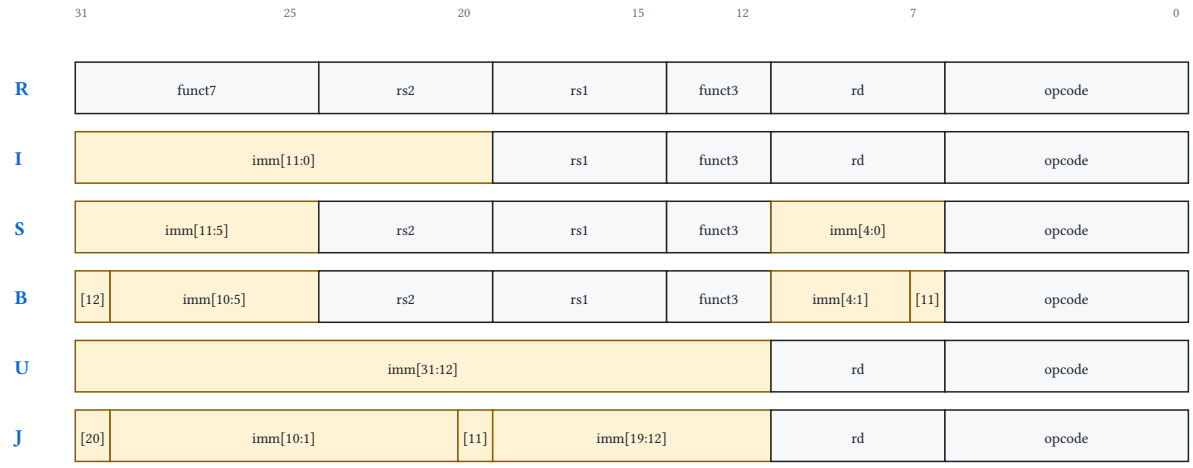
LEARNING GOALS

- Explain why the opcode and register fields sit at fixed positions.
- Explain the immediate layout in terms of multiplexer inputs.
- Avoid the trap that makes every negative `addi` subtract.
- Test a decoder with legal cases adjacent to its illegal ones.

Fixed fields, decoded in parallel

Every other field's *meaning* depends on the opcode, so the opcode must be decodable before anything else — which is why it is always `inst[6:0]`. Putting it at a fixed position means the register fields, `funct3` and the immediate can all be extracted **in parallel** with decoding the opcode rather than after it.

The register fields never move either: `rs1` is always `inst[19:15]`, `rs2` always `inst[24:20]`, `rd` always `inst[11:7]`. So lesson 13 addresses the register file *before the instruction has been understood at all.*



The opcode never moves, and neither do rs1, rs2 or rd. Only the shaded immediate fields are rearranged.
Every format puts its SIGN BIT in bit 31 — so sign extension needs no multiplexer at all.
B and J look shuffled because bit positions were chosen to stay aligned ACROSS formats, not in order within one.

The six formats. Only the shaded immediate fields are rearranged; every other field stays where it is.

The immediates look like vandalism

B puts `imm[11]` in `inst[7]` while `imm[12]` is in `inst[31]`. J shuffles its bits out of order entirely. Read as a table it looks arbitrary.

Two rules explain all of it:

1. **The sign bit is always `inst[31]`.** Every format. So sign extension needs no multiplexer at all — just replicate one fixed wire.
2. **Each immediate bit comes from as few different places as possible.** `imm[10:5]` is `inst[30:25]` in I, S, B *and* J. `imm[4:1]` is `inst[24:21]` in I and J, `inst[11:8]` in S and B.

Rule 2 is *why* B and J look shuffled: bit positions were assigned to stay aligned **across** formats, not in order **within** one.

MEASURED count the multiplexer inputs and the design appears
Of the 32 immediate output bits, counting how many distinct instruction-word positions feed each — a bit with *N* sources needs an *N*-input multiplexer:

sources	bits
1 — no multiplexer at all	**24**
2	7
3	1

41 multiplexer inputs across all 32 bits, against an upper bound of 160 if the five layouts were uncorrelated — 26% of it.

And the single bit needing three sources is `imm[11]` — precisely the one that looks most arbitrary in the specification. The designers put the ugliness in *one place* rather than spreading it everywhere.

WHY THIS MATTERS
"The encoding is weird" and "the encoding is carefully optimised" turn out to be the same observation from two sides. You cannot tell which from reading the specification; you can tell immediately from the multiplexer count.

Note also that B and J immediates are implicitly multiplied by two — bit 0 is always zero and is not stored, because instructions are 2-byte aligned. One extra bit of branch range, free.

The trap

Lesson 10's ALU takes `{inst[30], funct3}`, which falls straight out of the instruction word for R-type. For I-type it very nearly does:

```
addi x1, x2, -1
```

has `imm = 0xFFF`, so `inst[31:20]` is all ones and **`inst[30]` is 1**. Pass `{inst[30], funct3} = {1, 000}` to the ALU and it performs **SUB** instead of ADD — so `addi` with any negative immediate silently subtracts.

`inst[30]` is only meaningful for `OP-IMM` when `funct3 == 101`, distinguishing `SRLI` from `SRAI`. Everywhere else in `OP-IMM` it is part of the immediate and must be forced to zero.

WATCH OUT
This is invisible in an R-type-only test, and any immediate at or beyond 1024 — or any negative one — trips it. Without the fix, lesson 13's first program computes nonsense.

A BUG THIS COURSE ACTUALLY HIT
A bug in the first draft of this decoder, caught by the linter rather than by a test: `alu_op` was assigned from `{alu_bit30, funct3}`, and a *separate* wire applied the "loads and stores must add" override — and was never connected to anything. Every load would have asked the ALU for whatever `funct3` happened to mean. One assignment per output.

Testing a decoder

The vectors are **real encodings**, generated by a small assembler and each labelled with the assembly it came from — every format, the sign-extended extremes of each, and eight malformed instructions that must raise `illegal`.

The legal contrast cases matter as much as the illegal ones. A decoder that rejected everything would pass all the illegal tests, so encodings sitting right next to them — `blt`, `bge`, `sra` — are asserted *not* to be flagged.

WATCH OUT
One of the lesson's own illegal vectors was simply wrong. `0x00314863` was written intending "branch with `funct3=011`"; it actually has `funct3=100`, which is `blt` and perfectly legal. The decoder correctly declined to

flag it and the test failed. Hand-computing instruction encodings is error-prone — compute them with a script and print the fields back to check.

A BUG THIS COURSE ACTUALLY HIT

Seven mutations; six caught immediately. The survivor forced S-type `imm[0]` to zero — and was invisible because **every store offset in the vectors happened to be even**. Adding `sb x3,7(x2)` catches it at once.

Unlike lesson 11's surviving mutation, which was benign because the design was redundant, this one was a *real missing test*. Both outcomes are useful, and telling them apart is the skill.

What it cost

MEASURED decode is cheap; execution is where the silicon goes

`immgen`	70 LUT4 + 27 muxes
`decode`	44 LUT4 + 12 muxes

The immediate generator costs more than the entire control decoder — consistent with the multiplexer count. Both are tiny beside lesson 10's ALU at 885 LUT4.

Lesson 13 — A single-cycle RV32I processor

CONCEPT	Fetch, decode, execute, memory and writeback in one clock — the simplest thing that is genuinely a CPU.
YOU BUILD	A core that runs real RISC-V assembly, verified by the program's own output.
NEEDS	Lesson 12. A board, and it is worth it.

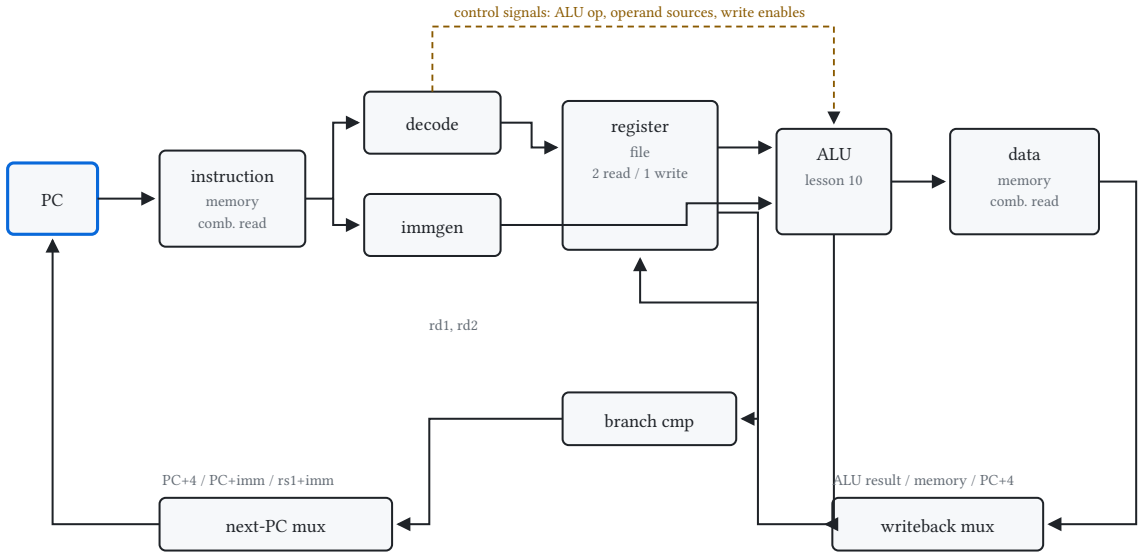
LEARNING GOALS

- See a CPU as a state machine whose state is a program counter.
- Explain why a single-cycle core needs separate instruction and data memories.
- Handle byte and halfword access, and sign versus zero extension.
- Test a processor through its architectural interface.
- Explain why the whole core runs at half the ALU's clock.

Nothing here is new

The ALU is lesson 10. The register file is lesson 11. The decoder and immediate generator are lesson 12. The only new logic is a branch comparator, the load/store aligners, and the next-PC multiplexer.

> **WHY THIS MATTERS**
> **A CPU is not a new kind of thing.** It is a state machine (lesson 05) whose state is a program counter, wired to parts you have already built and tested separately. That is the most demystifying idea in the course, and it is why those four lessons came first.



Everything shown happens in ONE clock period, so the critical path is this whole chain in series.
That is why the core closes at 61 MHz where the ALU alone managed 116 — and why lesson 16 pipelines it.

The datapath. Everything shown happens within one clock period.

Why single-cycle first

There are no pipeline registers between the stages, so every wire goes where the diagram says it goes. Read the source top to bottom and it is the picture, in order. Lesson 16 pipelines this exact core, and the comparison is the point of doing it this way round.

Harvard, and not as a preference

There are **two** memories, instruction and data, both read combinationally. That is not a style choice: a single-cycle CPU must fetch an instruction *and* access data in the same cycle, and one memory with one port cannot do both.

Real machines present one address space and manage with a single memory by pipelining — fetch in one stage, memory access in a later one — or by putting separate caches over shared memory. Here the simplest thing that works is two memories, and saying so is better than pretending it was taste.

The new parts

A separate branch comparator

Branches need six comparisons. The ALU could do them, but in a single-cycle CPU **the ALU is busy computing the branch target in the same cycle**. So a separate comparator, which is cheap and keeps the ALU's output free.

Note what RV32I omits: there is no `BGT` or `BLE`. `bgt a, b, L` assembles as `blt b, a, L` — the *assembler* swaps the operands and the hardware needs no extra cases. Six conditions, not ten.

Byte and halfword access

Memory is 32 bits wide; the ISA has `LB`, `LH`, `LW`, `LBU`, `LHU`, `SB`, `SH`, `SW`. So a byte store must change *one* byte and leave the other three alone — which is why the memory has per-byte write enables — and a byte load must select and extend. `LB` sign-extends and `LBU` zero-extends; both exist because C's `char` may be either and a compiler must be able to ask for exactly one.

The bug only running software could find

The CPU ran on the first attempt — 109 cycles, reaching `halt` — and got most answers right. Every failure was an **X**, and all of them traced to `LUI`.

A BUG THIS COURSE ACTUALLY HIT

Lesson 12's decoder carried a comment saying *"rs1 is x0 by convention for LUI"*, and the datapath relied on it. **That is false.** `LUI` has no `rs1` field at all: `inst[19:15]` is part of its 20-bit immediate. So the ALU was adding the immediate to whatever register those five bits happened to name — an undefined register — and the X propagated into every value built with a large constant.

Lesson 12's own testbench checked that `LUI` asks the ALU for ADD, and it does. **It never checked where the operands came from.** 161 passing decoder checks and a mutation-tested testbench did not find this; eighty-eight instructions of real software found it in one go.

WHY THIS MATTERS

Unit tests verify the parts. Integration finds the assumptions between the parts. Both are necessary and neither substitutes for the other. The decoder now has an explicit "operand A is zero" output, and its testbench asserts it for `LUI` and asserts it is *not* set for `add` or `auipc`.

Testing through the architectural interface

The test program stores every result to a known memory address, and the testbench reads data memory rather than peeking at the register file.

WHY THIS MATTERS

That tests the CPU through the interface *software actually uses*, so a core that computes correctly but *stores* wrongly still fails — which is also all that software could tell you. The program ends by setting a flag, and writes a marker if any branch goes the wrong way; the test requires the flag and forbids the marker.

What it cost, and what it cost in time

MEASURED thirty-two flip-flops in an entire CPU

LUT4	1,765

muxes	773
`CCU2C`	80
flip-flops	**32** — the program counter, and nothing else
`DPR16X4`	32 — the register file, as distributed RAM

All the state a single-cycle machine has is its PC. The register file lives in distributed RAM and everything else is combinational. Lesson 16 will add hundreds of flip-flops, and that is precisely what pipelining *is*.

MEASURED the number this lesson exists to produce

32-bit ALU alone (lesson 10)	115.94 MHz
the whole single-cycle core	**61.24 MHz**

Adding fetch, register read, memory access and writeback to the ALU's path **roughly halves the clock**. Every instruction — even `nop` — pays the full critical path, because every instruction takes exactly one cycle and that cycle must be long enough for the slowest one.

WHY THIS MATTERS
That is the entire argument for pipelining, now with a number attached. Lesson 16 splits this path into five stages so each is short, and pays for it in hazard logic.

On the board

Set the switches to 2 and you are reading the result of `addi t2, t0, -1`, which must be **9**. Had lesson 12's trap gone unfixed it would read 11.

This is a RISC-V processor you built, running a program you wrote, on a chip you configured. Worth a moment.

Lesson 14 — Formal verification

CONCEPT	Proving a property holds for all inputs up to a bound, instead of testing the cases you thought of.
YOU BUILD	A `riscv-formal` proof that the lesson 13 core really implements RV32I.
NEEDS	Lesson 13.

LEARNING GOALS

- Explain what a bounded model checker does, and what "bounded" excludes.
- Write assertions, assumptions and covers — and recognise when an assumption has made a proof vacuous.
- Read a counterexample trace as a minimal failing program.
- Judge when testing is enough and when it is not.

Why this lesson exists here

Lesson 13 ended with twenty-four passing checks and a CPU that ran. It also ended by finding a bug that **161 passing decoder checks and a mutation-tested testbench had missed** — because the bug was in an assumption no test had asked about.

> **WHY THIS MATTERS**
> That is the argument for formal methods in one sentence. A testbench tells you a design works for the inputs you imagined. For a processor that is nowhere near good enough: the bugs live in combinations nobody imagines, and the space is far too large to sample.

The idea

A bounded model checker takes a property — "this signal is never high while that one is low" — and either proves it holds for every reachable state within k cycles of reset, or hands you a concrete trace where it fails.

The word **bounded** is doing real work. A proof to depth 20 says nothing about a bug that needs 21 cycles to manifest. Depth is traded against solver time, and choosing it is part of the craft.

Three kinds of statement:

`assert`	something that must always be true. The thing being proved.
`assume`	a constraint on the inputs. **The dangerous one**: assume too much and you prove something about a machine nobody will build.
`cover`	a request for a trace reaching some state. The antidote to the above — if a `cover` is unreachable, your assumptions have excluded reality.

`riscv-formal`

An instruction-level specification of RV32I that an implementation can be checked against, by Claire Wolf. The core exposes a trace interface reporting what it retired each cycle — which instruction, which registers it read and wrote, what values — and the framework checks each retirement against the specification.

What it proves, instruction by instruction: that the result matches the ISA, that the program counter advanced correctly, that no register was disturbed that should not have been, and that memory access was properly aligned.

> **WHY THIS MATTERS**
> This is also where you find out which of the earlier lessons were quietly wrong. A formal proof over a CPU exercises the ALU, the register file and the decoder far harder than their own testbenches did — and it does so without anyone having thought of the case.

What it actually did

> **MEASURED** **the ALU, proved rather than sampled**

Lesson 10 computed that checking all $2^{32} \times 2^{32} \times 10$ ALU inputs would take **1,460 years** at a billion a second, and settled for 1,440 directed vectors plus 20,000 random ones.

The proof covers every one of them in **one second**. The ALU is combinational, so depth 1 examines the entire input space — for combinational logic, "bounded" costs nothing at all.

The register file, which has state, is proved for **every sequence of twelve operations** over every address and every data value, in 58 seconds.

The register file proof uses the technique worth taking away:

```
(* anyconst *) reg [4:0] watched;   // the SOLVER picks the address
always @(posedge clk)
    if (we && rd == watched && watched != 0) shadow <= wdata;
always @(*)
    if (shadow_valid && rs1 == watched) assert (rd1 == shadow);
```

`anyconst` makes this a proof rather than a spot check: the solver chooses `watched` and must satisfy the property for *every* choice, so one shadow register establishes a fact about all thirty-two.

Catching a real bug

The lesson ships a deliberately broken ALU — lesson 10's, with signed less-than using the shortcut that omits the overflow term. That failing region is a vanishing fraction of 2^{64} pairs, and 20,000 random vectors would very likely miss it.

MEASURED the solver found it in under a second

```
a = 0x80008000   signed  -2147450880
b = 0x57FE8001   signed   1476296705
signed:  -2147450880 < 1476296705  →  1
broken ALU answered:  0
```

Signs differ, the difference overflows, and the shortcut is wrong — precisely the case lesson 10 described. Nobody told the solver to look there.

WATCH OUT

The first thing a formal tool finds is usually a bug in your *property*. The first run of this proof failed with the design perfectly correct: the property used `y[sh-1:0]`, which is not a legal variable part-select, against a nonsensical right-hand side. Expect this.

The failure mode that makes formal dangerous

A BUG THIS COURSE ACTUALLY HIT

A proof can pass and be worth nothing. The lesson's `vacuous.v` uses the **real** ALU, makes **true** assertions, and **passes** — because of one line:

```
assume (a == b);
```

With the operands forced equal, `a < b` is false for every input the solver may consider. The assertion becomes `0 == 0`, the overflow logic is never exercised, and an ALU returning constant zero for SLT would pass identically.

Nobody writes that deliberately. They write an assumption to exclude an input the design is not meant to handle, it is slightly too strong, and the hole is invisible — because a passing proof looks the same either way.

The antidote is `cover`, and it is not optional. Running the vacuous example's cover task reports both covers **unreached**, which is the only signal that the proof above was hollow.

WHY THIS MATTERS

> **A formal proof is only as good as its assumptions, and nothing in a passing result tells you whether they were right.** The real ALU proof carries seven covers for this reason — that signed and unsigned comparison genuinely disagree, that `SRA` differs from `SRL`, that addition wraps — and all seven are reached.

What this lesson does not do

`riscv-formal` itself — checking a whole core against an instruction-level specification of RV32I — needs the core to expose a retirement trace interface and the riscv-formal repository alongside. That is a lesson's work on its own and is not here yet. What *is* here is the technique it rests on, applied to the two components the CPU is mostly built from.

Formal verification is not a replacement for testing. It is bounded, and writing a property that is both true and useful is a skill. What it does that testing cannot is **search** rather than sample.

Lesson 15 — A system on chip

NOT YET BUILT
The scope, ordering and theory of this lesson are settled, and the chapter below is the plan. The code and the measured results are not here yet – lessons are built in order, and this book reports only numbers that were actually measured. Nothing in this chapter is a measurement.

CONCEPT	A memory map: making peripherals look like addresses, so software can reach them.
YOU BUILD	Your CPU running C compiled by `riscv-gcc`, printing over the UART and drawing on the OLED.
NEEDS	Lesson 14, a board, and the GNU RISC-V toolchain.

LEARNING GOALS

- Decode an address space into RAM, ROM and peripherals.
- Explain what `volatile` in C corresponds to in hardware.
- Build the toolchain path from C source to a hex file in block RAM.
- Explain what must happen before `main`, when there is no operating system.

Where a CPU becomes a computer

Lesson 13's memories ignore the high address bits entirely, so address `0x0000_0100` and address `0xDEAD_0100` hit the same word — the memory *aliases* across the whole space. That is fine for one program and wrong in general.

A real system **decodes** the high bits to choose between RAM, ROM and peripherals, and faults on addresses that match nothing. That decoding is the memory map, and it is the difference between a CPU and a computer.

A peripheral is then just an address that does something when written. The UART of lesson 09 becomes a transmit register at some address and a status register at another, and the software-side consequence is one C keyword:

```
volatile uint32_t *uart_tx = (uint32_t *)0x8000_0000;
```

WHY THIS MATTERS
`volatile` tells the compiler that this location can change without the program changing it, and that writes to it have effects beyond storing a value — so it may not cache, reorder or elide them. It is the C type system acknowledging that some addresses are hardware, and it exists *because* memory-mapped I/O does.

The toolchain path

```
riscv-gcc  →  objcopy  →  .hex  →  $readmemh  →  block RAM
```

Lesson 08 established that `$readmemh` runs at elaboration time, so the compiled program is baked into the bitstream. Changing the C means recompiling *and* resynthesising — which is the point at which a soft CPU's real advantage appears: a design where behaviour lives in software is a design you can change without re-running place and route, if you load the program at run time instead.

Before `main`

With no operating system, something must set up the stack pointer, zero the `.bss` section, copy initialised data into RAM, and only then call `main`. That is `crt0`, and writing it is the clearest possible explanation of what a C runtime actually does — because it is about fifteen instructions and you can read all of them.

The payoff

The point at which the FPGA stops being a logic exercise and becomes a computer you programmed in C — and the point at which the hardware/software boundary becomes something you personally defined, which is the clearest way there is to understand it.

Lesson 16 — Pipelining

> **NOT YET BUILT**
> The scope, ordering and theory of this lesson are settled, and the chapter below is the plan. The code and the measured results are not here yet – lessons are built in order, and this book reports only numbers that were actually measured. Nothing in this chapter is a measurement.

CONCEPT	Overlapping instructions, and the hazard logic that makes overlap safe.
YOU BUILD	A five-stage version of the same core, running the same programs, at a higher clock.
NEEDS	Lesson 15.

LEARNING GOALS

- Explain why pipelining raises the clock without making any stage faster.
- Identify data, control and structural hazards.
- Explain which hazard forwarding can fix and which it cannot.
- Measure cycles-per-instruction and say where the cost went.

The argument, already made

Lesson 10 measured a 32-bit ALU at 115.94 MHz. Lesson 13 measured the whole single-cycle core at **61.24 MHz** — the ALU plus fetch, register read, memory and writeback, all in series, all paid for by every instruction including `nop`.

Pipelining does not make any of those steps faster. It overlaps them: while one instruction is in the ALU, the next is reading registers and the one after is being fetched. The clock period then only has to cover the *longest single stage* rather than the whole chain.

IF	instruction fetch
ID	instruction decode and register read
EX	execute — the ALU
MEM	data memory access
WB	write back to the register file

Between each pair sits a pipeline register. **This is where lesson 13's thirty-two flip-flops become hundreds**, and that growth is precisely what pipelining is.

The cost: hazards

Overlap is only safe if the overlapping instructions do not interfere, and mostly they do.

Data hazards. An instruction needs a result that an earlier one has computed but not yet written back. The fix is *forwarding*: route the result from where it already exists — the EX or MEM stage output — directly to where it is needed, rather than waiting for the register file.

The one forwarding cannot fix. A load's result is not available until the end of MEM, so an instruction immediately after a load that uses its result must *stall* — one bubble, every time. Compilers know this and schedule around it, which is one of the places where an ISA's design and a compiler's behaviour meet.

Control hazards. By the time a branch is resolved, the two instructions behind it have already been fetched. Either flush them — paying two cycles on every taken branch — or predict, and pay only when the prediction is wrong.

What to measure

Three numbers make the trade concrete:

1. the new maximum clock, against lesson 13's 61.24 MHz
2. the cycles-per-instruction on a real program, which is now above 1 because of stalls and flushes
3. the resource cost of the hazard logic and pipeline registers

WHY THIS MATTERS
Pipelining is the idea that turned processors from simple to subtle. Doing it to a core you already have, with a proof you already wrote in lesson 14, is the honest way to see what the complexity buys and what it costs — because the specification has not changed, so the lesson 14 proof should still pass.

Lesson 17 — Extensions

NOT YET BUILT
The scope, ordering and theory of this lesson are settled, and the chapter below is the plan. The code and the measured results are not here yet – lessons are built in order, and this book reports only numbers that were actually measured. Nothing in this chapter is a measurement.

CONCEPT	Privileged state, traps, and the extensions that make a core usable rather than merely correct.
YOU BUILD	Control and status registers, a machine-timer interrupt, and hardware multiply.
NEEDS	Lesson 16.

LEARNING GOALS

- Explain what a trap is and what hardware must save before taking one.
- Distinguish an exception from an interrupt.
- Write an interrupt handler in C and know what the hardware did first.
- Decide when to spend a hardware multiplier and when to iterate.

Optional, and genuinely so

The course's goal is met at lesson 16. This chapter exists because "a correct RV32I core" and "a core that can run an operating system" are separated by exactly these pieces, and it is worth seeing how few they are.

Traps

A trap is a forced, controlled change of control flow. The hardware must save enough to resume: where it was (`mepc`), why it stopped (`mcause`), and where to go (`mtvec`). `MRET` reverses it.

exception	caused by the instruction being executed — an illegal instruction, a misaligned access. Synchronous, and the instruction did not complete.
interrupt	caused by something outside — a timer, a peripheral. Asynchronous, and taken between instructions.

WHY THIS MATTERS
Lesson 12's decoder has produced an `illegal` output since the beginning, and lesson 13's core halts on it rather than executing nonsense. That signal cost one wire and has been waiting for this lesson, which is the point: designing the error path in from the start is far cheaper than retrofitting it.

The machine timer

`mtime` counts, `mtimecmp` holds a threshold, and an interrupt fires when one reaches the other. That is the whole mechanism, and it is enough to build a pre-emptive scheduler on — which is how a timer interrupt becomes an operating system.

The M extension

`MUL` maps onto the hardware multipliers lesson 08 met by accident — the part has 156 of them. `DIV` has no hardware at all and becomes a multi-cycle state machine, which is why division is slow on nearly every processor ever built.

WHY THIS MATTERS
By now the question "should this be hardware or software?" has a method rather than an opinion attached. Measure what the hardware costs, measure what the software costs in cycles, and decide — which is what this whole course has been teaching.

Appendix A — Board pinout

The authoritative constraint file is `common/ulx3s_v20.lpf`. The names below are the port names to use in a top-level module; the constraint file binds them to ECP5 balls. Icestudio uses the same friendly names on its canvas, so a block bound to `btn_1` and a Verilog port `btn[1]` land on the same ball.

Clock

`clk`	G2	25 MHz, free-running

Every lesson in this course is clocked by this and nothing else — one clock domain, no gated clocks, except lesson 04 which is about the exceptions.

LEDs

Active high: write 1 to light. `led_0` is the least significant bit of the `led[7:0]` port used throughout.

`led_7`	`led_6`	`led_5`	`led_4`	`led_3`	`led_2`	`led_1`	`led_0`
H3	E1	E2	D1	D2	C1	C2	B2

Buttons

`btn_0` is the power button and is **inverted** — it reads 1 when released. The rest are active high with pull-downs. All of them need debouncing, which is lesson 03.

`btn_0`	D6	power — inverted
`btn_1`	R1	FIRE1
`btn_2`	T1	FIRE2
`btn_3`	R18	UP
`btn_4`	V1	DOWN
`btn_5`	U1	LEFT
`btn_6`	H16	RIGHT

Slide switches

Static configuration input; no debouncing needed in practice.

`sw_0`	`sw_1`	`sw_2`	`sw_3`
E8	D8	D7	E7

OLED — SSD1331, 96 × 64, 16-bit colour

Four-wire SPI, write-only. `oled_dc` selects command (0) or data (1). Maximum clock about 6.6 MHz, which is why lesson 06 divides the 25 MHz clock by four.

`oled_clk`	P4	SPI clock, driven by the FPGA
`oled_mosi`	P3	data, most significant bit first
`oled_dc`	P1	command (0) or data (1) — a level, not a stream bit
`oled_resn`	P2	reset, active low; needs a real delay after release
`oled_csn`	N2	chip select, active low

USB serial — FTDI FT231X

> **WATCH OUT**
> Named from the **FTDI's** point of view. The FPGA *drives* `ftdi_rxd` and *reads* `ftdi_txd`, so `output ftdi_rxd` and `input ftdi_txd` are correct and look backwards. `ftdi_txden` must be driven high to enable the FPGA's driver. This is the classic place to lose an afternoon.

`ftdi_rxd`	L4	the FTDI receives — the FPGA **transmits**
`ftdi_txd`	M1	the FTDI transmits — the FPGA **receives**
`ftdi_txden`	L3	drive high to enable
`ftdi_nrts`	M3	
`ftdi_ndtr`	N1	

Everything else

The board definition lists 184 pins in total: 32 MB of SDRAM (`sdram_*`, 16 bits wide), digital video out (`gpdi_*`), the audio resistor ladders (`audio_*`), the SD card (`sd_*`), SPI flash (`flash_*`), an ADC (`adc_*`), the ESP32 interface (`wifi_*`), and 28 differential GPIO pairs (`gp0–gp27` and `gn0–gn27`). See `common/ulx3s_v20.lpf` for the complete set.

Part facts worth remembering

FPGA	Lattice ECP5 `LFE5U-85F-6BG381C`
Package	381-ball caBGA
Speed grade	6
LUT4s	83,640
Flip-flops	83,640
Block RAM	208 × 18 Kbit, about 3.7 Mbit
18×18 multipliers	156
PLLs	4
Clock	25 MHz
Programmer	FTDI FT231X — `openFPGALoader -b ulx3s`

Appendix B — Command reference

Per-lesson verbs

From inside a lesson directory, or from the project root with `L=<number>`:

verb	tool	what it does
`sim`	iverilog + vvp	compile source and testbench, run it, print PASS or FAIL
`wave`	GTKWave	open the VCD this lesson produced
`lint`	Verilator `-Wall`	width mismatches, inferred latches, multiple drivers
`graph`	yosys + graphviz	render the design as an SVG schematic
`synth`	yosys `synth_ecp5`	resource report: LUT4s, flip-flops, block RAM, carry
`bit`	yosys → nextpnr → ecppack	a bitstream, with timing closure
`prog`	openFPGALoader	load into FPGA SRAM — volatile, for iterating
`flash`	openFPGALoader `-f`	write to SPI flash — survives power-off
`clean`		delete build products

From the project root

```
make list              # which lesson is which number
make deploy L=0        # build lesson 00 and load it onto the board
make sim    L=5        # simulate lesson 05
make wave   L=3        # open lesson 03's waveform
```

`make check`	verify the toolchain and smoke-test lesson 00
`make sim-all`	every lesson's testbench — the regression suite
`make lint-all`	Verilator across every lesson
`make graph-all`	render every lesson as a schematic
`make ice`	regenerate the Icestudio projects from their specs
`make tasks`	regenerate the editor's per-lesson task list
`make docs`	build the HTML version of the lessons
`make pdf`	build this book
`make clean-all`	remove build products everywhere

Driving the tools directly

```
# simulate
iverilog -g2005 -Wall -o tb.vvp design.v tb_design.v && vvp tb.vvp

# strict lint
verilator --lint-only -Wall --top-module design design.v

# resource usage
yosys -p "read_verilog design.v; synth_ecp5 -top design; stat"

# resource usage at a different parameter value
yosys -p "read_verilog design.v; chparam -set W 32 design; \
          synth_ecp5 -top design; stat"

# schematic
yosys -p "read_verilog design.v; proc; opt_clean; \
          show -format dot -prefix d design"
```

```
dot -Tsvg d.dot -o d.svg

# place and route for the ULX3S 85F
nextpnr-ecp5 --85k --package CABGA381 --freq 25 \
  --json design.json --lpf common/ulx3s_v20.lpf --textcfg design.config
ecppack design.config design.bit

# program
openFPGALoader -b ulx3s design.bit        # SRAM
openFPGALoader -b ulx3s -f design.bit     # SPI flash
openFPGALoader --detect -b ulx3s          # is the board there?

# PLL parameters
ecppll -i 25 -o 100 --clkout1 40 -f pll.v
```

Reading a resource report

```
   LUT4              18     combinational logic, 4 inputs each
   TRELLIS_FF        26     flip-flops
   CCU2C             13     carry chain cells — TWO BITS EACH
   DP16KD             1     an 18 Kbit dual-port block RAM
   TRELLIS_DPR16X4    8     distributed RAM: 16 words x 4 bits, from LUTs
   PFUMX / L6MUX21    4     wide multiplexers, for functions over 4 inputs
   MULT18X18D         0     a hardware multiplier
```

The 85F has 83,640 LUTs, 83,640 flip-flops and 208 block RAMs, so a design using 34 cells is using about 0.04% of the part. Knowing roughly what things cost is the point: a 32-bit adder is 17 carry cells, a 32-bit barrel shifter is about 667 LUTs, and a whole RISC-V CPU is about 1,765.

Which simulator when

iverilog	the default. Fast to start, forgiving, good `$display`.
Verilator	much faster for long runs, but compiles to C++ and is strict about things iverilog tolerates. Used here mainly as a linter.
sby	not a simulator at all. Proves properties over *all* inputs up to a bound, instead of the handful your testbench thought of. Lesson 14.

Tools written for this course

`tools/install.sh`	set up a machine; no sudo, nothing system-wide
`tools/check.sh`	verify every tool runs, then smoke-test lesson 00
`tools/code.sh`	launch VS Code with the toolchain on `PATH`
`tools/rvasm.py`	an RV32I assembler — labels, pseudo-instructions, `--list`
`tools/check-asm.sh`	cross-validate that assembler against GNU `as`
`tools/mkice.py`	generate Icestudio projects from lesson Verilog
`tools/mktasks.py`	regenerate the editor's per-lesson tasks
`tools/mkwave.sh`	render WaveDrom timing diagrams to SVG
`tools/svg-flatten.py`	resolve CSS variables so diagrams print
`tools/mkdocs.sh`	build the HTML lessons
`tools/mkpdf.sh`	build this book

Appendix C — Verilog style rules

Conventions every lesson follows. Each exists because it corresponds to a mistake that is easy to make and annoying to debug.

Synchronous design, one clock

Everything is clocked by `clk_25mhz`. No gated clocks, no latches, no asynchronous resets unless a lesson is explicitly about them.

If you need something slower, divide with a counter and use the result as an **enable**, not as a clock:

```
// Yes: one clock domain, a slow enable.
reg [15:0] div = 0;
wire tick = (div == 16'd49999);
always @(posedge clk) div   <= tick ? 16'd0 : div + 1'b1;
always @(posedge clk) if (tick) count <= count + 1'b1;

// No: a derived clock. Does not use the clock network, will not meet
// timing predictably, and makes the design unanalysable.
always @(posedge div[15]) count <= count + 1'b1;
```

Blocking versus non-blocking

- `<=` inside `always @(posedge clk)` — always, without exception.
- `=` inside `always @(*)` combinational blocks.

Non-blocking assignments all take effect at the end of the time step, which is what a bank of real flip-flops does. Blocking assignments take effect immediately, which chains them.

```
// Shift register: each stage sees the OLD value of the previous one.
always @(posedge clk) begin
    a <= in;  b <= a;  c <= b;      // three flip-flops
end

// With blocking assignments this collapses: a, b and c all become `in`.
always @(posedge clk) begin
    a = in;   b = a;   c = b;       // ONE flip-flop, replicated
end
```

Lesson 03 measures it: three flip-flops against one.

Do not infer latches

A combinational `always` block must assign every output on every path, or synthesis infers a latch — storage with no clock, which is almost never what you wanted.

```
// Latch: `y` keeps its old value when sel is 2'b11.
always @(*) case (sel)
    2'b00: y = a;  2'b01: y = b;  2'b10: y = c;
endcase

// Fixed: assign a default FIRST, then override.
always @(*) begin
    y = 1'b0;
    case (sel)
        2'b00: y = a;  2'b01: y = b;  2'b10: y = c;
        default: y = 1'b0;
    endcase
end
```

> **WATCH OUT**
> In a *clocked* block the same omission is correct. `if (en) q ⩽ d;` with no `else` is a flip-flop with an enable, not a latch, because a flip-flop already holds its value between edges. Same syntax, opposite meaning, decided by whether the block is sensitive to a clock edge or to its inputs.

Widths

State widths explicitly and keep them consistent. Verilog will silently zero-extend, truncate or sign-extend, and the result is a bug that appears only at a boundary value.

```
reg  [7:0] count;
count ⩽ count + 1'b1;        // explicit 1-bit literal, not bare `1`
wire [7:0] sum   = a + b;    // drops the carry — deliberate? say so
wire [8:0] sum_c = a + b;    // keeps it
```

Two traps worth naming, both silent:

Shifts	`a << b` is as wide as `a`. So `1'b1 << sel` shifts a one-bit value and is always zero. Widen the literal before shifting.
Carries	`a + b` into a W-bit destination throws the carry away. Widen the intermediate to W+1.

And when a truncation is intended, say so — widen, then slice:

```
localparam [31:0]   LAST32 = DEPTH - 1;
localparam [AW-1:0] LAST   = LAST32[AW-1:0];
```

Reset and initialisation

The ECP5 initialises flip-flops from the bitstream, so `reg x = 0;` works and is used throughout this course. That is a convenience of this family, not a portable habit — on an ASIC you need an explicit reset. Where a lesson needs a real reset, it is synchronous:

```
always @(posedge clk)
    if (rst) count ⩽ 0;
    else     count ⩽ count + 1'b1;
```

One driver per signal

A register is driven from exactly one `always` block, and priority between its inputs is expressed with `if`/`else`. Two blocks assigning one register have no defined ordering in simulation and force synthesis to guess — Verilator rejects it, correctly.

Naming

`clk_25mhz`	the board clock, named for its frequency
`*_n`, `*n`	active low — `oled_resn`, `csn`
`*_i`	an internal copy of an output you also need to read
`S_*`	state constants, as `localparam`
`tb_*`	the testbench for the module of the same name
`*_next`	the combinational next-state value of a register

Name states as constants, never bare numbers in a `case`. `2'd2` tells a reader nothing; `S_GREEN` tells them everything.

Testbenches

Every testbench ends with an explicit verdict:

```
if (errors == 0) $display("PASS: %0d vectors checked", n);
else             $display("FAIL: %0d of %0d wrong", errors, n);
```

`make sim-all` greps for `PASS`, so this is the regression suite's interface rather than decoration.

Three further rules, each learned the hard way in a lesson:

1. **Compute the reference a different way from the design.** A model that shares the design's mistake proves nothing. Lesson 10's shifts are checked against an independent iterative implementation for exactly this reason.
2. **Set up monitors before stimulus.** Lesson 09's bit-period measurement asserted `start` and then waited for an edge that had already passed. Use `fork`/`join`.
3. **Do not assert what you have not established.** Lesson 09's testbench claimed mistimed frames were never silently corrupted, and passed. The claim was false. A test that asserts something false and passes is worse than no test.

Mutation testing

A testbench that has never caught a bug is unproven. Break the design deliberately and check the test notices. A surviving mutation means one of two things, and distinguishing them is the skill:

the test is incomplete	lesson 12: every store offset in the vectors happened to be even, so a bug in the immediate's low bit was invisible
the code is redundant	lesson 11: `x0` is enforced on both the read and the write path, so removing either alone changes no behaviour

Appendix D — Where to go next

Specifications worth having open

The RISC-V ISA manual, volume 1	Unprivileged instructions. Free, readable, and short enough to actually read — which is unusual for an ISA specification and is part of why this course can end where it does.
The RISC-V ISA manual, volume 2	Privileged architecture: the CSRs, traps and the machine timer that lesson 17 is about.
Lattice ECP5 family datasheet	The authoritative numbers for the part: setup and hold times, block RAM configurations, PLL ranges.
SSD1331 datasheet	The display controller of lessons 06 and 07. Worth reading once as an example of the genre.

Books

Patterson & Hennessy, *Computer Organization and Design*, RISC-V edition	The standard text, and the one this course's lessons 10 through 16 run parallel to. Where this book measures, that one explains — they complement each other well.
Harris & Harris, *Digital Design and Computer Architecture*	Starts further back and covers the Verilog alongside the architecture. Closest in spirit to this course.
Nisan & Schocken, *The Elements of Computing Systems*	The Nand2Tetris book. Builds a complete machine from NAND upward in simulation, and then a compiler and an operating system on top of it. A different and excellent answer to the same question.

Projects to read

picorv32	A small, well-documented RV32I core by Claire Wolf. The natural next thing to read after finishing lesson 13 — a production-quality version of what you just built.
SERV	The world's smallest RISC-V core, by Olof Kindgren — bit-serial, around 200 LUTs. A superb demonstration that "correct" and "fast" are independent axes.
VexRiscv	A configurable RV32 core generator. Shows what a pipeline, caches and an MMU look like at production scale.
LiteX	An SoC builder. Where you go when you want a CPU *and* SDRAM, Ethernet, video and an SD card, without building all of them yourself.
The ULX3S examples	`github.com/emard/ulx3s` — SDRAM controllers, video output, emulators of older machines. Written for this exact board.

Things this course deliberately left out

Each would make a good next project:

1. **An SDRAM controller.** The board has 32 MB of it, 16 bits wide, and it needs refresh, row activation and careful timing. The single hardest thing on this board, and the thing that turns a toy into a computer.
2. **Digital video output.** Serialisation at ten times the pixel clock, which needs the PLL of lesson 04 used in anger.

3. **Caches.** Lesson 16′s pipeline makes memory the bottleneck; a cache is the answer, and cache coherence is where it gets genuinely hard.
4. **A compiler back end.** Lesson 15 uses GCC. Targeting your own instruction set with LLVM is a large project and an education in why ISAs look the way they do.
5. **Formal proof of the pipeline.** Lesson 14 proves the single-cycle core. Doing the same for lesson 16′s pipelined version is strictly harder and strictly more valuable, because that is where the subtle bugs live.

A closing thought

The most transferable thing in this course is not Verilog, and it is not RISC-V. It is the habit of asking *what did that actually cost* and then running the two-second command that answers it.

Four times in these lessons, a rule that everyone knows turned out to be wrong for the fabric in front of us — and in each case the measurement took less time than arguing about it would have. That habit works on tools you have not met, on chips that do not exist yet, and on problems that are not about hardware at all.